

Estimating the State-Transition Model for Sequential Campaign Allocation

The Political-State Analog of a Volatility Process, and the Value of Waiting It Implies

Callum Doty

July 2026

Contents

1	Introduction	5
1.1	The Question This Paper Answers, and the One It Does Not	5
1.2	The Finance Analogy, Made Precise	5
1.3	What Motivates This Now	6
2	Related Literature	7
2.1	Option Pricing and Real Options	7
2.2	State-Space Estimation and Time-Series Processes	7
2.3	Political Science: Strategic Response and Forecasting	7
2.4	Research Gap	7
3	Problem Formulation	9
3.1	Notation	9
3.2	The State Vector $\mathbf{X}_t$	9
3.3	Why Calibration Precedes Optimization	10
3.4	The Value Function and the Stopping Problem	10
4	Theoretical Framework	12
4.1	Opponent Reaction $R_t(D_t)$	12
4.1.1	Why This Is Not “Just Another Stochastic Term”	12
4.1.2	Specification	12
4.2	The National Environment Process G_t	12
4.2.1	Candidate Specifications	12
4.2.2	λ for the Idiosyncratic-Uncertainty Decay Proxy	13
4.3	Race-Level Idiosyncratic Uncertainty $\varepsilon_{i,t}$	13
4.3.1	What Already Exists, and What Is Missing	13
4.3.2	A Data Constraint, Likely Permanent, and Why the Obvious Fallback Is Dangerous	13
4.3.3	Proposed Treatment	13
4.4	The Deploy-Branch Convolution Identity	14
4.5	From a Calibrated P to Θ : Standard Machinery, One Necessary Adaptation	15
4.6	A Scope Boundary That Must Be Stated Before Estimating Anything	16

5	Data	17
5.1	Data Sources	17
5.2	Coverage	17
5.3	Feature Engineering	17
5.4	Cleaning	17
5.5	Final Dataset	18
6	Parameter Estimation & Calibration	19
6.1	Opponent Reaction $\hat{\eta}(\text{tier})$	19
6.2	The National Environment Process $\sigma_G(\Delta t)$	20
6.3	Idiosyncratic Uncertainty Decay	21
6.4	Candidate-Spending Trickle Rate	21
6.5	Calibration Summary	22
7	Optimization Algorithm	23
7.1	Optimization Strategy	23
7.2	Algorithm	23
7.3	Computational Complexity	23
7.4	Computational Environment	23
8	Empirical Results	25
8.1	Internal Validation	25
8.2	A Necessary Condition: Does the Optimal Policy Actually Depend on Information? .	26
8.3	Main Result Under the LP Control Approximation – Status Suspended (See Section 8.9)	27
8.4	A Comparative-Static Horizon Extension (Not a Historical Counterfactual)	28
8.5	Testing the Binary Framing Directly: A Continuous Deployment Fraction	29
8.6	The Correction That Produced This Result	31
8.7	Mechanism Decomposition: Isolating Information Value from Deterministic Sequencing	32
8.8	Statistical Rigor: Simulation Noise and Out-of-Sample Policy Evaluation	33
8.9	A Critical, Unresolved Allocator-Robustness Finding	34
9	Discussion	40
9.1	Why the Corner Flipped (Within the LP Approximation – Itself Now Suspended) . .	40
9.2	What Remains Structurally Unpriced – and What the Decomposition Shows Is Already Priced	40
9.3	Reconciling with Paper II’s Baseline	41
9.4	A Domain-Intuition Check Against Real Data	41
9.5	Generalizability	41
10	Limitations	42
10.1	Allocator Robustness (Critical, Unresolved – See Section 8.9)	42
10.2	Data	42
10.3	Modeling	42
10.4	Computational	42
10.5	Practical Deployment	43
11	Conclusion	44

12 Data Availability	45
13 Code Availability	45
14 Conflict of Interest	45
15 References	46
16 Appendix A: Notation Reference Table	47
17 Appendix B: Derivations	47
17.1 B.1 The Deploy-Branch Terminal-Probability Identity (Corrected)	47
17.2 B.2 Incremental Decomposition of the Idiosyncratic-Uncertainty Schedule	47
17.3 B.3 EMA/OU Term-Structure Fit	48
18 Appendix C: Proofs	48
18.1 C.1 $\Theta(T) = 0$ (Boundary Condition)	48
19 Appendix D: Additional Robustness Analyses	48
19.1 D.1 Leave-One-Cycle-Out Stability Check	48
19.2 D.2 Necessary-Condition Check Detail	48
20 Appendix E: Sensitivity Analyses	49
20.1 E.1 LP-vs-Nonlinear Allocator Divergence in the Deploy Branch	49
20.2 E.2 Sensitivity of $\Theta(0)$ to the Codebase-Wide Corrections in Paper I	49
21 Appendix F: Hyperparameters	49
22 Appendix G: Pseudocode	50
23 Appendix H: Database Schema	50
24 Appendix I: Additional Figures	50
25 Appendix J: Additional Tables	50
26 Appendix K: Configuration Files	51
27 Appendix L: Reproducibility Checklist	51

Abstract

A companion paper (Paper II) builds a sequential architecture for deploying campaign capital – a rollout policy that re-solves a single-period optimizer at each reporting period, distinct from full model-predictive control – and discovers a gap this policy cannot express: capital held rather than spent has option value, because retaining it preserves the ability to react to information that has not yet arrived. It names this gap Θ , states the Bellman equation that would price it, and stops there deliberately, because that equation contains an expectation operator with nothing specified to integrate over. This paper argues that treating the Bellman equation as the remaining problem is backward, in the same sense that asking for the value of a stock option is backward before anyone has specified how the stock price moves: an option’s value is a corollary of a specified stochastic process for the underlying, not an independently derived quantity.

This paper is the political-science analog of specifying that process. We define the campaign state vector explicitly, decompose its evolution into a fully-derived control component (inherited from Paper I) and three components requiring genuine estimation – opponent reaction to a committee’s own spending, a national political-environment process, and race-level idiosyncratic uncertainty – and estimate each from public data where the data allows. Opponent reaction is estimated from a seven-cycle (2012–2024) panel of dated independent-expenditure filings, tiered by competitiveness, with a cycle-weighted (random-effects) point estimate of $\hat{\eta} \approx 0.26\text{--}0.38$ in genuinely contested tiers after confirming most tiers show real cycle-to-cycle variation rather than a single stable constant. The national environment process is calibrated from a five-cycle historical generic-ballot series, giving a term structure $\sigma_G(\Delta t) \approx 0.18\text{--}0.20 \times \sqrt{\Delta t}$ that is empirically indistinguishable between a random walk and a mean-reverting process over the three-to-nine-month horizon that matters near Election Day. Race-level idiosyncratic uncertainty cannot be estimated as a genuine time-varying process under the public-data-only constraint this research program maintains, given that well under 10% of competitive House districts receive repeated public polling in a typical cycle; it is instead treated as an explicit, bounded proxy that borrows its resolution rate from the calibrated national process, a choice we defend and flag rather than present as equivalent to a fitted model. The assembled simulator is validated against held-out 2022 and 2024 data before being trusted for pricing (a September information set rank-correlates with eventual November outcomes at $\rho = 0.47$ and 0.65 respectively) and passes a four-part internal self-consistency check. Once specified and estimated, Θ is a standard application of regression-based Monte Carlo (Longstaff–Schwartz) backward induction, not a new contribution. Solved against the live 2026 decision using the fast LP within-period allocator, $\Theta(0)$ is substantially positive – $+4.5$ to $+5.1$ expected seats across three calibration scenarios at a 98-day horizon. ****This LP-based figure is currently suspended as a policy conclusion.**** A preliminary allocator-robustness check – replacing the LP allocator with the true, diminishing-returns-respecting nonlinear allocator throughout the entire backward induction, not merely at the immediate decision – reverses the result: under identical simulated opponent-response draws, $\Theta_{\text{LP}}(0) = +5.097$ becomes $\Theta_{\text{nonlinear}}(0) = -0.376$ at a reduced path count ($K = 15$). That path count is too small to establish the sign of $\Theta(0)$ under the correctly specified allocator with confidence, but the roughly 5.5-seat gap is far too large to be ordinary Monte Carlo variation around the LP figure. The defensible conclusion at this stage is narrower than this paper’s own headline claim: the estimated value of waiting is highly sensitive to the within-period control solver, and the LP-based estimate cannot presently be interpreted as a policy-robust value of waiting. Statements to the effect that holding is the unanimous recommendation, or that waiting is robustly worth $+4.6$ to $+5.9$ seats across calibration scenarios, are results of the LP approximation, not yet demonstrated properties of the underlying nonlinear model, and this paper suspends its practical recommendation pending a higher-powered nonlinear evaluation. This reversal was found through the same discipline the paper’s central methodological argument insists on – and the mechanism decomposition, statistical-rigor checks, and two prior corrections reported below (a missing spending-drift channel, then a variance-specification double-count) remain valid diagnostics of the LP-based calculation; what they diagnose can no longer be presented as an estimate of $\Theta(0)$ itself. This paper’s contribution is the specified and estimated transition law, P , and the discovery that the value it implies for Θ is not yet identified under the model’s own intended optimization; resolving that is the immediate next step of this research line, not a footnote to it.

Keywords: optimal stopping, real options, state-space estimation, regression-based Monte Carlo, Ornstein–Uhlenbeck process, campaign finance, stochastic control

1 Introduction

1.1 The Question This Paper Answers, and the One It Does Not

Paper II asks how a committee should deploy capital sequentially, given a valuation model, and in building that architecture discovers a gap: nothing in it prices the difference between “this race is worth funding” and “this race is worth funding *right now*.” Paper II calls the missing object Θ , states the Bellman equation it would need to solve, and stops there deliberately.

The natural next move is to treat that Bellman equation as the remaining problem and try to solve it. This paper argues that is the wrong next move, for a precise reason: a Bellman equation over an unspecified state-transition law is not merely difficult – it is not yet a well-posed mathematical object. Consider the equation itself,

$$V_t(\mathbf{X}_t, F_t) = \max_{0 \leq \mathbf{p}_t \leq F_t} \mathbb{E}_P[V_{t+1}(\mathbf{X}_{t+1}, F_t - \mathbf{1}'\mathbf{p}_t + \text{fundraising}_t)]$$

Every symbol here is defined except one: P , the law governing how $\mathbf{X}_{t+1}$ is generated from $\mathbf{X}_t$. Without P , this expression has an expectation operator with nothing inside it to take an expectation *over*. It is not that solving this equation is hard without P – it is that the equation does not yet say anything without P . Specifying P is therefore not a preliminary step before the real work; it is the entire remaining scientific content of the problem. Once P exists, computing V_t (and therefore Θ) is standard: simulate forward paths under P , and use a regression-based Monte Carlo method (Longstaff and Schwartz 2001) to estimate continuation values by backward induction. That machinery is decades old and well understood. It is not what makes this problem hard. **This paper’s job is to specify and, where the data allows, estimate P .** Everything about Θ , reserve policy, and optimal stopping is a downstream implication of that specification, not a separate derivation.

1.2 The Finance Analogy, Made Precise

The correspondence to option pricing is exact enough to be load-bearing, not decorative.

Table 1: The Option-Pricing Correspondence

Finance	This project
Stock price S_t	Campaign state $\mathbf{X}_t$ (Section 3.2)
Price process $dS_t = \mu S_t dt + \sigma S_t dW_t$	State-transition law P (Sections 4.2–4.4)
Volatility σ	G_t ’s innovation variance, $\varepsilon_{i,t}$ ’s decay structure, opponent reaction $R_t(D_t)$
Option value $V(S_t, t)$	Value function $V_t(\mathbf{X}_t, F_t)$
Early-exercise value / American option premium	$\Theta(t)$, the value of not yet committing capital
Black–Scholes PDE / binomial tree / Longstaff–Schwartz	The Bellman recursion above, solved once P is known

No one asks “what is the value of waiting to exercise this option?” before specifying σ . The question is not answerable in that order. The same is true here: Paper II’s Bellman equation, exactly like an option-pricing PDE, is the *consumer* of a specified process, not a substitute for specifying one.

1.3 What Motivates This Now

Paper II’s live 2026 run gives this urgency beyond theoretical motivation. At approximately four months from Election Day, Paper II’s rollout policy – which implicitly assumes P contributes nothing beyond what the current period already knows – recommends a majority of deployable capital to Likely R and Safe R seats, including a nontrivial allocation to deeply-safe districts (Paper I’s persuasion-ceiling-corrected pipeline: 51.4% of a live \$393M reserve). This is the empirical size of the gap a specified P closes, and Section 8 reports what closing it changes about the live recommendation.

2 Related Literature

2.1 Option Pricing and Real Options

The mathematical machinery this paper ultimately applies – regression-based Monte Carlo pricing of an American-style option with early-exercise features – is due to Longstaff and Schwartz (2001), building on the foundational option-pricing framework of Black and Scholes (1973) and Merton (1973). The extension of that machinery from financial to real (physical or organizational) assets is the subject of real-options theory (McDonald and Siegel 1986; Dixit and Pindyck 1994), which this paper’s companion (Paper II) already introduces as the conceptual frame for uncommitted campaign capital. This paper’s relationship to that literature is specific: the machinery is not this paper’s contribution (Section 1.1), and is applied here essentially unmodified once P exists (Section 4.6).

2.2 State-Space Estimation and Time-Series Processes

Specifying P requires standard time-series tools applied to a genuinely novel domain. The choice between a random walk and a mean-reverting (Ornstein–Uhlenbeck) specification for a bounded political quantity is a classical model-selection problem (the Ornstein–Uhlenbeck process itself: Uhlenbeck and Ornstein 1930); realized-volatility term-structure estimation, used here to calibrate the national-environment process’s horizon-dependent standard deviation directly from data rather than assuming a parametric form, follows the general logic of realized-volatility estimation in financial econometrics (Andersen and Bollerslev 1998). State-space filtering more broadly (Kalman 1960) is the natural refinement of the smoothing rule Paper II uses for its own state update, though this paper does not implement a full Kalman treatment (Section 10.2).

2.3 Political Science: Strategic Response and Forecasting

Paper I already situates the campaign finance and election-forecasting literatures relative to this research program; two threads bear directly on this paper specifically. Erikson and Palfrey (2000) model campaign spending as a simultaneous strategic game, establishing the theoretical prior this paper’s opponent-reaction estimate (Section 4.2, 6.1) tests directly with data: that a committee’s own spending decision changes, rather than simply adding to, the opposing side’s. Gelman and King (1993), already cited in Paper I for the fundamentals-versus-polls decomposition motivating a race’s expected margin as a structural quantity distinct from its moment-to-moment noise, is the same distinction this paper’s separation of $\mu_{i,t}$ (control-driven) from $\varepsilon_{i,t}$ (idiosyncratic shock) formalizes for an allocation, rather than forecasting, purpose.

2.4 Research Gap

No existing work specifies and estimates a full state-transition law for a political campaign’s decision-relevant state, suitable for pricing the value of delaying an irreversible spending commitment. The option-pricing and real-options literatures supply the valuation machinery in general form; they do not supply a calibrated instance for this domain. The strategic-interaction and forecasting literatures supply pieces of the transition law – evidence that spending elicits a reaction, methods

for tracking a national environment factor – but neither assembles them into a joint process an optimal-stopping calculation can consume, and neither confronts the practical data constraints (Section 5) that determine which components of that process can actually be estimated from public data and which cannot.

3 Problem Formulation

3.1 Notation

Table 2: Notation (in addition to Papers I and II)

Symbol	Meaning
P	the state-transition law governing $\mathbf{X}_{t+1} \mid \mathbf{X}_t$ (this paper’s subject)
G_t	national generic-ballot point estimate (D – R) at period t
$\varepsilon_{i,t}$	race-level idiosyncratic shock to $\mu_{i,t}$
$R_t(D_t)$	opponent (Republican-aligned) spending as a reactive function of the committee’s own control
$\eta(\text{tier})$	opponent-reaction coefficient, tiered by Cook competitiveness
$\sigma_G(\Delta t)$	the national-environment process’s realized-volatility term structure
λ	decay rate governing how race-level idiosyncratic uncertainty resolves over time
$V_i(t)$	cumulative remaining idiosyncratic-uncertainty target at time t for race i
$V_t(\mathbf{X}_t, F_t)$	the value function: expected seats achievable from state $\mathbf{X}_t$ with deployable capital F_t remaining
$\Theta(t)$	$V_t^{\text{wait}} - V_t^{\text{deploy-now}}$, the value of waiting
K	number of simulated Monte Carlo paths

3.2 The State Vector $\mathbf{X}_t$

Paper II leaves $\mathbf{X}_t$ informal and its transition operator f fully generic. Before any component of P can be specified, $\mathbf{X}_t$ itself needs a precise definition. Grounding this in the sequential architecture’s own implementation, the state at reporting period t is

$$\mathbf{X}_t = \left(\{\mu_{i,t}, \sigma_{i,t}, D_{i,t}, R_{i,t}, L_{i,t}\}_{i=1}^N, G_t, F_t \right)$$

where, per race i : $\mu_{i,t}$ and $\sigma_{i,t}$ are the smoothed expected-margin and margin-uncertainty estimates (Paper II’s valuation, re-estimated and EMA-smoothed each period); $D_{i,t}$ and $R_{i,t}$ are cumulative Democratic- and Republican-aligned spending to date; $L_{i,t}$ is committed capital already irreversible for that race. At the aggregate level, G_t is the national generic-ballot point estimate and $F_t = B_t - \sum_i L_{i,t}$ is deployable capital (Paper II’s ledger identity).

Table 3: Evolution Mechanism by Component

Component	Evolution mechanism	Status
$D_{i,t}$	Control: $D_{i,t+1} = D_{i,t} + p_{i,t}$, the committee’s own decision	Fully specified (the decision variable)
$\mu_{i,t}$	Control + two stochastic shocks: $\mu_{i,t+1} = \mu_{i,t} + \Delta\mu_i(p_t) + \beta_i\Delta G_t + \varepsilon_{i,t}$	Control term derived (Paper I); shocks require Sections 4.3–4.4
G_t	Stochastic process , exogenous to any single committee’s decisions	Requires Section 4.3
$R_{i,t}$	Reactive process , a function of the committee’s own control	Requires Section 4.2
$F_t, L_{i,t}$	Deterministic bookkeeping given B_t and the ledger identity	Fully derived (Paper II)
$\sigma_{i,t}$	Smoothed via Paper II’s EMA; no independent process proposed here	Explicitly out of scope

Of six rows, two are already fully solved by Papers I–II, one is definitional bookkeeping, and three are genuinely unresolved. This paper treats those three as co-equal objects of estimation, not as a primary process with two footnotes.

3.3 Why Calibration Precedes Optimization

This is stated as its own claim because it is this paper’s central methodological point: **without a calibrated P , the Bellman equation above is not difficult to solve – it is undefined.** “Difficult” describes a well-posed problem that resists an easy method. An expectation operator with no specified distribution to integrate against is not that; it is a placeholder. It would be a mistake to pick a convenient distribution for G_t and $\varepsilon_{i,t}$ purely so the equation can be exercised; that produces a number, but not a defensible one, and the number would carry false precision. The correct sequencing is: estimate P from data, report honestly where it cannot yet be estimated, and only then compute Θ – carrying forward, explicitly, whatever uncertainty a component of P that cannot be fully estimated implies (Section 4.4).

3.4 The Value Function and the Stopping Problem

Given a specified P , the value function satisfies the Bellman recursion of Section 1.1, with boundary condition $V_T(\cdot, F_T) = \text{Paper I’s static payoff at Election Day } (\Theta(T) = 0: \text{nothing left to wait for once the election has occurred})$. The object this paper ultimately computes is $\Theta(t) = V_t^{\text{wait}} - V_t^{\text{deploy-now}}$: the value of retaining F_t uncommitted for one more period, against the value of deploying it now

via Paper I's static optimizer, is a standard implication of P , not a new derivation, once P exists (Section 4.6).

4 Theoretical Framework

This section is the paper’s heart: specification of each stochastic component of P , the resulting fully-assembled state-transition model, and the machinery – standard once P exists – that turns it into Θ .

4.1 Opponent Reaction $R_t(D_t)$

4.1.1 Why This Is Not “Just Another Stochastic Term”

The other two stochastic components $(G_t, \varepsilon_{i,t})$ enter $\mathbf{X}_t$ ’s transition additively, alongside the control term, without altering what the control term *means*. Opponent reaction is different in kind: if Republican-aligned spending responds to Democratic increments at some rate η , the *effective* control is not $\Delta\mu_i(p_t)$ as Paper I derives it holding R_i fixed – it is $\Delta\mu_i(p_t)$ evaluated against a moving $R_{i,t}$ that partially offsets the increment. Paper I’s optimizer already encodes this mechanically:

$$R_i(D_i) = R_{i,\text{base}} + \eta \cdot \max(0, \text{party}_i - \text{party}_{i,\text{obs}}), \quad \eta \in [0, 1]$$

At $\eta = 0$, Paper I’s static gradient is exactly correct. At $\eta \rightarrow 1$ (dollar-for-dollar matching), the log-ratio term driving the entire marginal-seat-gain chain rule collapses toward zero – the committee’s spending decision stops mechanically converting into margin movement at all. **This is not a refinement to the control term; it is a statement about how much of the control term is real.** A single scalar η is very likely mis-specified: the economically plausible prior is that opponents match aggressively in Toss-Ups, where a marginal dollar changes the outcome, and largely ignore spending in races that are not competitive for either side. The primary specification is therefore $\eta(\text{tier})$, fit separately by Cook competitiveness tier, not a pooled scalar.

4.1.2 Specification

$$\Delta R_{i,t}^{\text{IE}} = \eta(\text{tier}_i) \cdot \Delta D_{i,t-1}^{\text{IE}} + u_{i,t}$$

with race fixed effects applied via within-(cycle, district) demeaning, pooled across a reconstructed and cleaned panel of date-bucketed independent-expenditure filings, η interacted with Cook tier. Section 5.4 documents the data-cleaning steps (amendment resolution, duplicate-transaction and implausible-amount filtering) this specification requires and that a cycle-cumulative analysis does not surface.

4.2 The National Environment Process G_t

4.2.1 Candidate Specifications

Two natural candidates: a random walk, $\Delta G_t \sim N(0, \sigma_G^2 \Delta t)$, under which variance grows unboundedly with horizon; and a mean-reverting Ornstein–Uhlenbeck process, $dG_t = \kappa(\bar{G} - G_t)dt + \sigma_G dW_t$, under which variance saturates. These have materially different implications for Θ in principle – a random walk implies option value keeps growing the longer one waits, while mean reversion implies

most resolvable uncertainty is realized within a bounded window – but the generic ballot is bounded by underlying partisanship, so formal model selection would very likely favor OU over a random walk almost regardless of the data, while the mean-reversion speed implied by a realistic lookback window is typically slow enough that over the three-to-six-month horizon actually relevant to Θ near Election Day, the two specifications are nearly indistinguishable in their implications (confirmed empirically, Section 6.2). The primary deliverable of this section is therefore not a model-selection verdict between the two, but the empirical, non-parametric term structure $\sigma_G(\Delta t)$ itself – realized volatility as a direct function of horizon, estimated from the historical series without forcing a Gaussian process form onto it first.

4.2.2 λ for the Idiosyncratic-Uncertainty Decay Proxy

Fitting $\text{Var}(\Delta G)(t) = A(1 - e^{-t/\tau})$ to the pooled historical term structure (Section 6.2) yields a decay rate $\hat{\lambda} = 1/\hat{\tau}$, used directly in Section 4.4’s idiosyncratic-uncertainty proxy below – the assumption being that idiosyncratic information resolves at a rate comparable to national information, absent any race-specific data to say otherwise.

4.3 Race-Level Idiosyncratic Uncertainty $\varepsilon_{i,t}$

4.3.1 What Already Exists, and What Is Missing

Paper I’s σ_i model is a **cross-sectional** fit – how much residual uncertainty a race of a given type carries once, not how that uncertainty evolves or resolves over a cycle. A genuine $\varepsilon_{i,t}$ process requires knowing how much of a race’s idiosyncratic uncertainty resolves per unit time, and how that resolution correlates across races beyond what G_t already captures.

4.3.2 A Data Constraint, Likely Permanent, and Why the Obvious Fallback Is Dangerous

Estimating this genuinely requires district-level polling history across the competitive universe. In a typical recent cycle, well under 10% of House districts receive two or more public polls at all, and most of those cluster in the final weeks before Election Day – not enough to fit a genuine time-series process for $\varepsilon_{i,t}$ per race; for the large majority of races the effective sample size is one observation or zero. A tempting fallback – treating $\varepsilon_{i,t}$ as a static draw from Paper I’s cross-sectional σ_i distribution – is not a neutral simplification; it is wrong in a specific, consequential direction. Treating $\varepsilon_{i,t}$ as a static draw implicitly assumes all of a race’s idiosyncratic uncertainty resolves instantaneously, equivalent to assuming there is no idiosyncratic information left to arrive – exactly backward, since it would make Θ collapse toward zero for the wrong reason (nothing left to wait for) rather than the right one (nothing left to *learn*).

4.3.3 Proposed Treatment

We instead borrow the resolution *rate*, not the distribution, from Section 4.3’s calibrated national process, applied as a shrinkage/decay factor on the static cross-sectional σ_i :

$$\sigma_{i,t} = \sigma_i^{\text{static}} \cdot \sqrt{1 - e^{-\lambda(T-t)}}$$

where λ is fit from $\sigma_G(\Delta t)$'s term structure (Section 6.2) rather than estimated separately per race. This is explicitly a proxy, not a fitted $\varepsilon_{i,t}$ process, and is reported as one: it prevents Θ from being inflated by an implausible “surprise arrives all at once on Election Day” assumption, but it does not constitute having estimated race-level dynamics from data, because the data to do so does not exist in usable quantity (Section 10.1).

4.4 The Deploy-Branch Convolution Identity

A committee that deploys its full reserve at period t does not thereby learn nothing about how the future would have unfolded; it forecloses the ability to *act* on that information, but the terminal win probability the deploy branch should be scored against still reflects the true distribution of where μ_i ends up by Election Day, not merely its period- t value. Formally, if $\mu_i(T) = \mu_i(t) + \xi$ for a mean-zero remaining shock $\xi \sim N(0, V_i(t))$ (the movement not yet realized at t , i.e. the part of the idiosyncratic-uncertainty budget σ_i^2 that has not yet resolved as of t), then, applied to a $\mu_i(t)$ that is itself already a clean structural estimate with no resolved shock embedded in it,

$$\mathbb{E}_\xi[\Phi((\mu_i(t) + \xi)/\sigma_i)] = \Phi\left(\frac{\mu_i(t)}{\sqrt{\sigma_i^2 + V_i(t)}}\right)$$

a standard Gaussian-convolution identity (Appendix B.1 derives it). This lets the “deploy now” branch of the backward induction analytically integrate over the expected effect of all remaining idiosyncratic drift in one step, rather than requiring the simulator to wait through it path by path. The identity is exact for what it integrates over – the expected effect of unresolved, mean-zero Gaussian shocks on terminal win probability – and Section 8.4 shows it has a direct and important consequence for how Θ behaves as the horizon lengthens: it does not, however, capture the value of *adaptive* decision-making, that new information arriving mid-campaign could change *how* the reserve is allocated, not merely *whether* it is deployed (Section 9.2 returns to this scope boundary directly).

A critical asymmetry. The convolution identity above requires ξ to be genuinely mean-zero. This holds for idiosyncratic shocks and, absent any real spending growth while a committee waits, trivially held for $D_{i,t}$ itself in an earlier version of this model (Section 8.6's correction). It does *not* hold once $D_{i,t}$ is given a real, deterministic, non-zero-mean growth process (Section 6.4) – a committee's own candidate committees continue raising and spending money organically while the party committee waits. A convolution that omits this deterministic drift understates the deploy branch's true value, biasing the comparison toward “wait” for the wrong reason. Section 8.6 reports finding and correcting exactly this asymmetry.

A variance double-count, found and corrected in this same audit. The identity above is stated for a $\mu_i(t)$ with *no* resolved shock embedded in it – σ_i^2 and $V_i(t)$ are meant as two disjoint pieces of the same idiosyncratic-uncertainty budget (resolved-by- t and remaining-after- t),

not two independent, additive sources. The implemented simulator, however, constructs $\mu_i(t) = \mu_i^{\text{struct}}(t) + \varepsilon_i^{\text{cum}}(t)$, where $\varepsilon_i^{\text{cum}}(t)$ is the *already-realized* share of the idiosyncratic budget as of t (Appendix B.2’s telescoping construction: $\text{Var}(\varepsilon_i^{\text{cum}}(t)) + V_i(t)$ is constant in t). Applying the convolution above to this already-partly-resolved $\mu_i(t)$ and *still* adding the full σ_i^2 term prices the same idiosyncratic budget twice – once through the realized simulated shock, once again through the convolution’s own σ_i^2 term – inflating the deploy branch’s effective variance by up to a factor of roughly 2–3 at horizons well before Election Day (confirmed numerically: at the live 98-day horizon, $\text{Var}(\varepsilon_i^{\text{cum}}(t)) + V_i(t) = 91.97$ for a representative $\sigma_i = 15$ throughout, versus $\sigma_i^2 + V_i(t)$ ranging from 225 (near T) to 317 (at $t = 0$) under the uncorrected formula). The corrected convolution, given a $\mu_i(t)$ that already embeds the resolved share, uses $V_i(t)$ alone:

$$\Phi\left(\frac{\mu_i(t) + \Delta\mu_i}{\sqrt{V_i(t)}}\right)$$

with the boundary at $t = T$ ($V_i(T) = 0$) handled by the same limit that already anchors $\Theta(T) = 0$ (Appendix C.1) rather than by re-introducing σ_i there. This is a strictly more consequential correction than Section 8.6’s trickle-drift fix: it moves the live-horizon $\Theta(0)$ from the +1.3–+1.7 range reported in earlier drafts of this section to +4.5–+5.1 (Table 11), because the uncorrected formula suppressed *all* win probabilities toward 0.5 at every period before T , muting precisely the late-campaign certainty that “wait, then deploy once things resolve” depends on. External review subsequently caught that this same treatment had been left inconsistent specifically at the terminal boundary itself for the default (stochastic) path – the fix above, as now stated, is the *fully* corrected version, applied uniformly to every period including $t = T$; the intermediate-only version of the fix (reported in an earlier pass of this section) had moved the figures further, to +4.6–+5.9, before this last refinement. Section 8.6 reports all three corrections together, in the order they were found.

4.5 From a Calibrated P to Θ : Standard Machinery, One Necessary Adaptation

Once Sections 4.2–4.5 specify P , computing $V_t(\mathbf{X}_t, F_t)$ and therefore $\Theta(t)$ is standard: simulate forward paths of $\mathbf{X}_t$ under P ; at each step, regress simulated continuation values on a basis of current-state features (Longstaff and Schwartz 2001); proceed by backward induction from T . We state this explicitly so it is not mistaken for this paper’s contribution – the contribution is Sections 4.2–4.5; this section is what those sections are *for*.

The one adaptation this domain requires: compress the regression basis. The state vector carries per-race features for $N > 400$ races. Regressing simulated continuation values on a per-race feature basis is not merely slow – at the sample sizes involved ($K \sim 10^4$ paths, ~ 15 time steps, giving $\sim 10^5$ observations against a feature matrix with hundreds of columns) it is a data-analytic mistake that would overfit badly and say more about simulation noise than about Θ . The regression basis must instead be a small set of portfolio-level aggregate features, the same design used in basket-option pricing when the same dimensionality problem arises from pricing an option on many underlyings at once:

1. $\mathbb{E}[\text{Seats}]_t = \sum_i \Phi(\mu_{i,t}/\sigma_{i,t})$ – the portfolio’s current level.

2. $\text{Var}[\text{Seats}]_t = \mathbf{1}'\Sigma_t\mathbf{1}$ – the portfolio’s current risk (Paper I’s covariance model).
3. $\max_i \text{MSG}_{i,t}$ – the value of the single best marginal dollar available right now.
4. A near-threshold count – the number of races within a small margin of the majority-determining threshold, a direct proxy for how much of the portfolio’s outcome is still genuinely undecided.
5. G_t – the current national-environment level, entered directly as a fifth feature (never fed into μ_i ’s structural formula, per Section 4.3’s scope boundary below).

4.6 A Scope Boundary That Must Be Stated Before Estimating Anything

The existing point-in-time state-reconstruction machinery this paper’s estimation and validation work reuses documents a constraint this paper must respect: α_3 , Paper I’s margin-model generic-ballot coefficient, was estimated entirely from *between-cycle* variation – one GB value per historical cycle, identical across every race in that cycle. Its identifying variation has never been within-cycle. Wiring the calibrated ΔG shock into μ_i for a within-cycle simulation step would apply α_3 to an estimand it was never fit against. **This paper does not make that substitution.** The assembled simulator computes ΔG (Section 4.3) and η -driven $R_{i,t}$ movement (Section 4.2), but validates each against real held-out data on its own terms (Section 8.1), and enters G_t into the Longstaff–Schwartz regression basis as an independent feature (item 5 above) rather than folding it into μ_i under a model specification never estimated to support it. Re-estimating the margin model on a panel with within-cycle GB observations is a prerequisite to that stronger integration, and is left for future work (Section 10.2).

5 Data

5.1 Data Sources

Beyond every source Papers I and II already use, this paper requires three additional, specifically dated panels. **Dated independent-expenditure filings**, resolved to specific filing dates rather than cycle-cumulative totals (FEC Schedule E, via the comprehensive bulk export), pooled across all seven cycles this repository has usable data for (2012–2024), used to estimate opponent reaction (Section 6.1). **A historical, dated, multi-cycle generic-ballot series** (2018, 2020, 2022, 2024), recovered via a Wayback Machine snapshot of FiveThirtyEight’s discontinued live generic-ballot data feed, since no other multi-cycle dated source was available in this repository – the standard `generic_ballot_by_cycle.csv` carries exactly one static value per cycle, no dates at all, and RealClearPolitics’s bulk archive was found to be blocked by bot protection when checked directly. **A dated candidate-committee financial panel** (FEC API `/committee/{id}/reports/` endpoint, per-filing-period Form 3 reports), used to calibrate the candidate-spending trickle rate the deployment-branch convolution correction (Section 4.5) requires; this endpoint returns genuinely dated, per-period filings, unlike the bulk `weball` files this project’s other candidate-spending measures use, which report only a single cycle-cumulative total.

5.2 Coverage

Opponent-reaction estimation pools all seven cycles with usable dated IE data (2012–2024, $n = 57,763$ delta-panel rows in the full extension). The national-environment term structure pools four historical cycles (2018, 2020, 2022, 2024) plus the live 2026 series. The candidate-spending trickle rate is calibrated on the same seven-cycle panel as opponent reaction. Simulator validation (Section 8.1) uses the 2022 and 2024 cycles as held-out targets. The main Longstaff–Schwartz solve (Section 8.3) is run against the live, in-progress 2026 state at a 98-day horizon, and against a 364-day counterfactual horizon holding all other state fixed (Section 8.5).

5.3 Feature Engineering

The opponent-reaction panel is constructed at biweekly reporting periods (matching Paper II’s reporting cadence) from January through early November of each cycle, with race fixed effects applied via within-(cycle, district) demeaning rather than a dummy-column matrix, since pooling roughly 800 district-cycle observations across many biweekly periods makes an explicit dummy matrix computationally wasteful for no estimation benefit. The national-environment term structure is computed directly from the recovered daily series at fixed horizons (30 through 450 days) without imposing a parametric form first.

5.4 Cleaning

Two real, checked data-quality issues bear directly on the opponent-reaction estimate and are documented rather than silently absorbed. First, district attribution (`can_office_dis`) is blank in only 0.1–0.3% of House-general independent-expenditure rows in every cycle from 2012 through 2024 – a specific a priori concern about pre-2018 data that does not hold in this data. Second, and more

consequentially, blank expenditure dates (`exp_date`, needed for the date-bucketed reconstruction specifically) are *worse* in the recent cycles this paper’s headline estimates rely on most (33.3% in 2022, 28.4% in 2024) than in older cycles (0.0% in 2012), and 2022’s raw file separately contains a parsing-artifact corrupted row and a broader issue of duplicated transaction identifiers (26% of all 2022 IE rows share a duplicated `tran_id`). Both were confirmed not to contaminate any figure already reported in Papers I–II, which consume cycle-cumulative rather than date-bucketed totals, but a `tran_id`-deduplication and implausible-amount filtering step is a documented prerequisite specifically for this paper’s transaction-level reconstruction.

5.5 Final Dataset

Table 4: State-Transition Estimation Dataset Summary

Quantity	Value
Opponent-reaction panel, cycles	2012–2024 (7 cycles)
Opponent-reaction panel, observations	57,763 (delta-panel rows, full 7-cycle extension)
National-environment series, cycles	2018, 2020, 2022, 2024 (historical) + live 2026
National-environment series, daily rows per cycle	1,142–1,174
Candidate-spending trickle panel, cycles	2012–2024 (7 cycles, same extension as opponent reaction)
Simulator validation held-out cycles	2022, 2024
Monte Carlo paths per scenario (K)	2,000
Reporting periods (live horizon)	7 (biweekly, 98 days to Election Day)

6 Parameter Estimation & Calibration


6.1 Opponent Reaction $\hat{\eta}(\text{tier})$

The primary, cycle-weighted estimate pools all seven usable cycles via a random-effects (DerSimonian–Laird) meta-analytic combination, which gives each *election* a vote scaled by its own precision, rather than letting whichever cycle happened to generate the most transaction rows dominate a naive transaction-weighted pool.

Table 5: Opponent-Reaction Estimates by Tier

Tier	2-cycle pool (2022+2024)	7-cycle naive pool	7-cycle cycle-weighted (DL)	I^2
Toss-Up	0.475	0.422	0.277	73%
Lean D	0.259	0.637	0.277	75%
Lean R	0.304	0.572	0.304	0%
Likely R	0.405	0.562	0.363	85%
Likely D	−0.165	0.180	0.043	46%
Safe D	0.648	0.371	0.380	70%
Safe R	−0.144	0.584	0.258	89%

I^2 – the share of cross-cycle spread that is real variation rather than sampling noise – is 70–89% for five of seven tiers, confirming a finding Section 8.1’s validation makes precise: most tiers, including Toss-Up and Lean D, show statistically significant cycle-to-cycle variation once tested with adequate power, not a single stable structural constant. **Precision (a small standard error within a given pool) is not the same claim as temporal stability (the same true value across cycles), and most tiers fail the stability test.** The cycle-weighted (DL) column is the best available point estimate – it does not let a transaction-heavy cycle dominate, and it uses all seven cycles rather than an arbitrary two-of-seven subset – and gives $\hat{\eta} \approx 0.26$ – 0.38 in genuinely contested tiers (Toss-Up, Lean D, Lean R, Likely R), well short of a naive dollar-for-dollar prior ($\eta \approx 1$). Estimates for the Safe tiers and Likely D are retained for completeness but are identified from sparse, lumpy transaction activity (well under 20% of periods show any IE spending at all in these tiers) and should be read with substantially more caution.



6.2 The National Environment Process $\sigma_G(\Delta t)$

Table 6: Generic-Ballot Realized Volatility Term Structure

Horizon (days)	Pooled std(ΔG), 5 cycles	std/ $\sqrt{\text{days}}$, all 5	std/ $\sqrt{\text{days}}$, 4 historical only
30	1.09	0.199	0.186
60	1.54	0.199	0.186
90	1.89	0.199	0.183
180	2.70	0.201	0.186
270	3.29	0.200	0.178
365	3.81	0.199	0.160
450	4.17	0.197	0.128

Pooled across genuinely independent cycles, $\text{std}/\sqrt{\text{days}}$ is remarkably stable from 30 through 270 days (0.183–0.201) rather than steadily declining, indicating an apparent mean-reversion signal visible in any single-cycle series is a small-sample artifact rather than a real feature of the process; a visible decline only emerges past roughly 365 days, consistent with Section 4.3’s prediction that a random walk is a good approximation over the three-to-nine-month range that matters for Θ near Election Day, with the random-walk/mean-reversion distinction only mattering at horizons beyond about a year. The working calibration is $\sigma_G(\Delta t) \approx 0.18\text{--}0.20 \times \sqrt{\Delta t \text{ (days)}}$ (historical-cycles-only preferred for methodological cleanliness), giving $\sigma_G \approx 2.0$ points at the roughly 120-day horizon relevant to the live application. Fitting $\text{Var}(\Delta G)(t) = A(1 - e^{-t/\tau})$ to this term structure gives $\hat{\lambda} = 1/\hat{\tau} = 0.00536 \text{ day}^{-1}$ ($\hat{\tau} = 186.5$ days), the decay rate Section 4.4’s idiosyncratic-uncertainty proxy uses directly.




Figure 2: Generic-ballot realized-volatility term structure, pooled across five cycles.

6.3 Idiosyncratic Uncertainty Decay

With $\lambda = 0.00536 \text{ day}^{-1}$ fixed from Section 6.2, race-level uncertainty resolution follows $\sigma_{i,t} = \sigma_i^{\text{static}} \sqrt{1 - e^{-\lambda(T-t)}}$ directly, with no additional free parameters. Figure 3 shows the resulting decay schedule.




Figure 3: Idiosyncratic uncertainty decay schedule, borrowed rate from the national-environment process.

6.4 Candidate-Spending Trickle Rate

A per-tier candidate-committee spending trickle rate – the deterministic drift the deploy-branch convolution correction (Section 4.5) requires – is calibrated from the seven-cycle dated candidate-financial panel (Section 5.1) using the mean daily disbursement rate within each tier, not the median: the median is exactly \$0.00/day in every tier, an artifact of comparing FEC’s quarterly

filing cadence against this project’s biweekly reporting grid rather than evidence of no real growth. Extending this calibration from the original two-cycle (2022, 2024) panel to the full seven-cycle panel reduced every tier’s fitted rate by 16–43% (e.g., Toss-Up: \$12,725/day $\rightarrow$ \$7,193/day) – the same pattern already documented for $\hat{\eta}$ in Section 6.1 – and is the calibration used in every result reported in Section 8.

6.5 Calibration Summary

Table 7: State-Transition Parameter Calibration Summary

Parameter	Value	Method
$\hat{\eta}(\text{Toss-Up})$	0.277	7-cycle DerSimonian–Laird
$\hat{\eta}(\text{Lean D})$	0.277	7-cycle DerSimonian–Laird
$\hat{\eta}(\text{Lean R})$	0.304	7-cycle DerSimonian–Laird
$\hat{\eta}(\text{Likely R})$	0.363	7-cycle DerSimonian–Laird
$\sigma_G(\Delta t)$	$0.18\text{--}0.20 \times \sqrt{\Delta t}$	Pooled realized volatility, 5 cycles
λ (decay rate)	0.00536 day^{-1} ($\tau = 186.5$ days)	Fit to σ_G term structure
Candidate trickle rate	tier-specific, e.g. Toss-Up \$7,193/day	7-cycle mean daily rate
K (Monte Carlo paths)	2,000	–
Regression basis dimension	5 features	Section 4.6

7 Optimization Algorithm

7.1 Optimization Strategy

The backward induction is Longstaff–Schwartz regression-based Monte Carlo: simulate K forward paths of $\mathbf{X}_t$ under the specified P ; at each step moving backward from T , regress simulated continuation values on the five-feature basis of Section 4.6; compare the regression-estimated continuation (“wait”) value against the deploy branch’s closed-form convolution value (Section 4.5); the chosen action at each simulated state is whichever value is larger.

7.2 Algorithm

Algorithm 1: Regression-Based Monte Carlo for Θ

Input: P (calibrated: $\eta(\text{tier})$, $\sigma_G(\text{dt})$, λ , trickle rate),
live state X_0 , K simulated paths, N periods to Election Day
Output: $\Theta(t)$ for $t = 0, \dots, N$; frac_deploy_now at each period

1. Simulate K forward paths of X_t under P :
 - $D_{i,t}$: control (fixed at 0 further discretionary spend on the wait branch) + deterministic candidate trickle drift
 - G_t : standalone zero-drift random walk at $\sigma_G(\text{dt})$ [never fed into μ_i ; Section 4.7's scope boundary]
 - $R_{i,t}$: $\eta(\text{tier}) * \delta(D_{i,t})$ + residual noise from the fitted regression residuals
 - $\epsilon_{i,t}$: independent increments matched to the cumulative decay schedule $V_i(t)$ (Section 4.4)
2. For $t = N$ (Election Day) down to 0:
 - a. Deploy branch: $V_{\text{deploy}}(X_t) = \text{closed-form convolution}$ (Section 4.5), corrected for trickle drift asymmetry
 - b. Wait branch: regress simulated V_{t+1} on the 5-feature basis (Section 4.6); $V_{\text{wait}}(X_t) = \text{fitted continuation value}$
 - c. $\Theta(t) = V_{\text{wait}}(X_t) - V_{\text{deploy}}(X_t)$
 - d. Optimal action: deploy if $V_{\text{deploy}} > V_{\text{wait}}$, else wait
3. Report $\Theta(t)$, frac_deploy_now (share of the K paths choosing "deploy" at each t), and basis R^2 as a diagnostic

7.3 Computational Complexity

Each period requires $O(K \cdot N_{\text{races}})$ work to simulate one step forward and $O(K)$ work for the basis regression (five features, independent of the race universe size by construction – the entire point of Section 4.6’s compression). With $K = 2,000$ paths and 7 reporting periods, a full binary-framing solve completes in minutes on commodity hardware; the continuous deployment-fraction generalization (Section 8.5) multiplies this by the number of grid points (5 or 11), still completing well within an hour.

7.4 Computational Environment

Implemented in Python, reusing Paper I’s estimation and optimizer code directly for the deploy branch’s within-period allocation (`optimizer.allocator.optimize()`, the fast LP

path, used for Monte Carlo tractability – $K \times N \times (\text{scenarios})$ calls would take days against Paper I’s full nonlinear solver). `scripts/solve_bellman_lsm.py` implements the binary hold-or-deploy framing; `scripts/solve_bellman_lsm_continuous_phi.py` implements the continuous deployment-fraction generalization; `scripts/estimate_eta_reaction.py`, `scripts/estimate_gb_volatility.py`, and `scripts/estimate_candidate_spend_trickle.py` implement the three calibration stages of Section 6. As in Papers I–II, a fixed random seed ensures exact reproducibility of every Monte Carlo result reported below.

8 Empirical Results

8.1 Internal Validation

Simulator self-consistency (the gate before any result is trusted). A multi-period simulator built from independently-calibrated components is not automatically trustworthy just because each component was validated on its own; per-step implementation errors can compound silently across many steps. Four checks were run at 5,000 simulated paths per cycle (2022, 2024) before any Θ figure was computed from the assembled simulator:

Table 8: Simulator Self-Consistency Checks

Check	2022	2024	Verdict
A: simulated $\sigma_G(\Delta t)$ vs. target, at 28/56/112 days	ratios 1.001/1.002/0.998	ratios 0.996/0.993/0.990	Matches
B: $\hat{\eta}$ recovered from simulated paths vs. input	0.344 vs. input 0.321 (rel. error 7.2%)	0.159 vs. input 0.175 (rel. error 9.3%)	Matches – no sign/scale bug
C: cumulative ε variance vs. target (mean/max rel. error)	1.5%/3.7%	1.3%/4.2%	Matches (Monte Carlo noise at $K=5,000$)
D: simulated margin spread / target remaining uncertainty	0.98–1.02	0.98–1.02	Matches

All four pass. This confirms the simulator faithfully implements what Sections 4.2–4.4 calibrated; it does not re-validate the calibration itself, which the next check addresses.

Validation A: does a September information set correctly rank-order eventual outcomes? Race state was reconstructed at a real September 1 snapshot for 2022 and 2024, $\mu_{i,\text{Sept}}$ computed via Paper I’s unmodified pipeline, and Spearman-correlated against realized November margin, restricted to the competitive set:

Table 9: Validation A – September-to-November Rank Correlation

Cycle	n	ρ	p
2022	61	0.467	< 0.001
2024	53	0.650	< 0.001

Both positive and significant: a September information set correctly rank-orders competitive races’ eventual outcomes two-plus months out, clearing the bar for the underlying valuation chain any simulated path is built from to be worth simulating at all.

Validation B: does $\sigma_G(\Delta t)$ match what actually happened? Comparing the calibrated term structure against realized $|\Delta G|$ over the real September 1 $\rightarrow$ Election Day window:

Table 10: Validation B – Realized vs. Predicted National-Environment Movement

Cycle	Δt (days)	Realized ΔG	Predicted SD	z	Within 2 SD?
2018	66	−0.16	1.49	−0.11	Yes
2020	63	−0.08	1.46	−0.05	Yes
2022	68	−2.18	1.52	−1.44	Yes
2024	65	−1.95	1.48	−1.31	Yes

All four realized moves are within 2 standard deviations of the calibration, and all four moved in the same direction (toward Republicans) in the final two months – suggestive of a real, if variable-strength, late-cycle tendency, but not a reliable, tradeable signal at $n = 4$ (mean $z = -0.73$ across all four cycles). G_t is not fed into μ_i ’s structural formula regardless (Section 4.7), so this remains an observation about a channel the model does not structurally use, not a calibrated input; Section 10.2 returns to it as a limitation rather than a finding.




Figure 4: Realized generic-ballot movement, September to Election Day, against the calibrated term structure.

8.2 A Necessary Condition: Does the Optimal Policy Actually Depend on Information?

Before asking whether *preserving* the ability to reallocate has positive value, a logically prior question must be settled: does new information even change what an optimal allocator wants to do? Re-solving Paper I’s optimizer against real historical state snapshots at a 60-day-out and a 14-day-out point in the 2022 and 2024 cycles, holding the total party budget fixed and identical so

that only the information available differs, shows roughly one-third to one-half of the top twenty targeted races changing between the two snapshots (Jaccard overlap 0.54 and 0.67 respectively) – real, substantial turnover. Even before considering the value of sequential decision-making, the allocation policy is demonstrably state-dependent: adaptive reallocation is a meaningful decision problem in this domain, not a theoretical abstraction being modeled for its own sake. This does not, on its own, imply that *waiting* has positive value (Section 8.3 answers that question directly, and the two findings are not in tension, as Section 9.2 discusses).

8.3 Main Result Under the LP Control Approximation – Status Suspended (See Section 8.9)

The figures in this section, and Tables 11–13 throughout this paper, are estimates under the LP within-period control approximation, not yet demonstrated properties of this paper’s intended nonlinear allocation model. Section 8.9 reports a preliminary allocator-robustness check finding these figures do not survive replacing the LP allocator with the true nonlinear allocator throughout the backward induction, and suspends the practical “hold the reserve” recommendation pending further investigation. The material below is retained because it remains a valid, internally-consistent description of what the LP-based calculation produces, which is itself informative (Section 8.9 argues the LP’s own discontinuous, knapsack-like behavior may be manufacturing apparent option value) – but it should not be read as this paper’s settled empirical conclusion.

Solved against the live 2026 state at a 98-day horizon, across three calibration scenarios – two single-cycle opponent-reaction brackets (fit on 2022 and 2024 individually, bounding a range of historically plausible reaction strength) and one scenario drawing ($\hat{\eta}$, resid_std) per simulated path from the full seven-cycle empirical distribution:

Table 11: $\Theta(0)$ at the Live 98-Day Horizon (LP Control Approximation)

Scenario	$\Theta_{LP}(0)$ (expected seats)	frac_deploy_now
eta_fit_2022	+4.546	0.0%
eta_fit_2024	+5.086	0.0%
eta_bootstrap_all_cycles	+4.692	0.0%

$\Theta_{LP}(0)$ is substantially positive in every scenario, and holding the deployable reserve is the unanimous recommendation *under the LP allocator*. This is the result of a calculation that has now been through three rounds of correction: an earlier, incorrect version found the opposite sign (missing the candidate-spending trickle channel entirely); the version after that found a positive but much smaller $\Theta(0)$ (+1.3 to +1.7) due to a variance-specification double-count in the deploy-branch convolution; and external review of that fix caught a further, related inconsistency at the $t = T$ boundary specifically, corrected in Section 4.5 and reported in full in Section 8.6. This last correction moved the figures modestly (roughly ± 0.4 seats per scenario) rather than by the multiple-fold swing the earlier double-count fix produced. A fourth, more serious issue supersedes these: Section

8.9 shows the LP-vs-nonlinear allocator choice itself moves $\Theta(0)$ by more than any of the three corrections above, in a direction that reverses its sign.




Figure 5: The binary hold-vs-deploy decision at $t = 0$: value of waiting against value of immediate deployment, across scenarios.

Table 11 reports only the $t = 0$ snapshot; the full per-period schedule underlying it is worth seeing directly rather than only in prose. Panel A below shows $\Theta(t)$ declining smoothly toward the $\Theta(T) = 0$ boundary (Appendix C.1) as Election Day approaches, in every scenario; Panel B shows the share of simulated paths for which deploying immediately is already optimal at that period, which stays near zero until roughly the final six weeks and rises only late in the cycle.




Figure 6: $\Theta(t)$'s full schedule across the live 98-day horizon, all three scenarios.

8.4 A Comparative-Static Horizon Extension (Not a Historical Counterfactual)

Holding today's actual state exactly fixed and re-running the identical backward induction with only the “days to Election Day” parameter changed from 98 to 364 tests whether the live-horizon

result is sensitive to how far $t = 0$ sits from Election Day. This is a comparative-static exercise, not a realistic historical counterfactual: it holds the current 98-day state’s polling, spending, and generic-ballot levels fixed and varies only the remaining-horizon parameter, which no real point in the 2026 cycle actually satisfies simultaneously (a race 364 days from Election Day would not, in reality, already carry today’s spend levels and polling position):

Table 12: $\Theta_{\text{LP}}(0)$ at the 364-Day Counterfactual Horizon (LP Control Approximation – Suspended, Section 8.9)

Scenario	$\Theta(0)$ (expected seats)	frac_deploy_now
eta_fit_2022	+4.093	0.0%
eta_fit_2024	+4.239	0.0%
eta_bootstrap_all_cycles	+4.713	0.1%

Under the corrected convolution, $\Theta(0)$ is smaller at the longer horizon in the two single-cycle brackets, and essentially unchanged (a 0.021-seat difference, within the Monte Carlo noise Section 8.8 quantifies) in the bootstrap scenario – the opposite direction, or a wash, from what an earlier draft of this section reported. That earlier claim (“hold is favored more, not less, the further out the decision is made”) was itself an artifact of the same variance double-count Section 4.5 corrects: the uncorrected formula kept re-adding a full σ_i^2 at every period regardless of how much time remained, so a longer horizon mechanically accumulated more of that spurious inflation. Once the double-count is removed, a longer horizon still favors hold in absolute terms (0.0–0.1% deploy-now at $t = 0$ in every scenario) – the qualitative conclusion is unchanged – but the corrected mechanism is that most of the option value is concentrated in the *final* few months rather than growing without bound the further out one stands, consistent with $V_i(t)$ ’s own shape (Section 4.4): remaining idiosyncratic uncertainty is bounded by σ_i^2 regardless of how far out t is, so extending the horizon past the point where most of that budget is already “in play” adds comparatively little. This is a genuine, reportable change in the mechanism’s implication, not a reversal of the paper’s headline recommendation.

8.5 Testing the Binary Framing Directly: A Continuous Deployment Fraction

The binary hold-or-deploy framing might be too coarse to express a genuine small-but-nonzero optimal reserve. This is tested directly rather than argued about, generalizing the backward induction to a genuine impulse-control problem over a discrete budget grid ($\{0, 0.25, 0.5, 0.75, 1.0\} \times F_0$, later confirmed at an 11-point, 10%-step grid), with unspent capital carried forward as a state variable rather than a one-time choice.

Table 13: Continuous Deployment-Fraction Generalization, 5-Point Grid (LP Control Approximation – Suspended, Section 8.9)

Scenario	$V(\text{hold})$	$V(\text{deploy})$	Gap (hold – deploy)	Chosen fraction
eta_fit_2022	236.845	233.610	+3.236 seats	100.0% hold
eta_fit_2024	237.472	233.721	+3.751 seats	100.0% hold
eta_bootstrap_all_cycles	236.032	233.310	+2.722 seats	67.1% hold, 23.2% at 25%, 0.9% at 50%, 8.9% at 75%, 0.0% full deploy




Figure 7: Continuous deployment-fraction generalization: value achieved at each grid point, by scenario.

The corner holds cleanly in the two single-cycle brackets, now by a wider margin than before correction. **eta_bootstrap_all_cycles** no longer places any mass at all on full immediate deployment (0.0%, versus 18.2% under the uncorrected convolution) and instead spreads its non-hold mass across the interior fractions – a genuinely different qualitative shape than the near-bimodal (0%/100%) pattern the uncorrected model produced, consistent with the corrected model no longer artificially compressing win probabilities toward 0.5 at intermediate budget levels. “Hold” still wins by a wide margin on average. (Note: which single-cycle bracket has the smallest gap has changed across each round of correction in this section – **eta_fit_2022** at one point, **eta_bootstrap_all_cycles** now – so this paper no longer singles one out as “closest to indifference”; the 11-point spot-check below covers both single-cycle brackets rather than picking whichever happens to be nearest at a given correction.) An 11-point spot-check confirms the 5-point result under the corrected convolution for both single-cycle brackets: **eta_fit_2024** (the representative case shown in the figure above’s panels A–B) gives gap = +3.762 seats versus +3.751 at 5 points (0.011-seat difference); **eta_fit_2022** gives gap = +3.244 versus +3.236 (0.008-seat difference). Both are 100.0% hold at every grid point from 10% through 90%, confirming the corner is not a grid-resolution artifact in either case.

8.6 The Correction That Produced This Result

Three corrections, found by checking mechanism rather than trusting output, separate the result reported above from two earlier, materially different findings.

A missing channel. An earlier version of the wait branch held candidate-committee spending $D_{i,t}$ fixed while waiting, because no per-filing-date source for candidate spending was believed to exist anywhere in this project’s data – itself later found to be incorrect (Section 5.1’s dated candidate-financial panel, recoverable from an FEC API endpoint not previously checked directly). With $D_{i,t}$ fixed, opponent reaction $\hat{\eta}$ had nothing to react to on the wait branch at all, and the resulting $\Theta(0)$ was negative (deploy favored) in every scenario tested. Once candidate committees’ real, dated spending growth is wired in as a genuine drift process, the wait branch legitimately captures value the deploy branch cannot: information about how each race’s own fundamentals continue developing while the party committee waits.

An asymmetric convolution, caught before trusting the first re-run. The first re-run under the corrected wait branch produced $\Theta(0)$ figures of +6.8 to +8.0 expected seats – an order of magnitude beyond anything else in this research line, which prompted checking the mechanism directly rather than reporting the number. The deploy branch’s “integrate over future drift in one step” convolution (Section 4.5) is only valid when the future movement it integrates over is mean-zero; with a real, deterministic, non-zero-mean trickle now driving $D_{i,t}$, the deploy branch was silently missing the expected μ appreciation from the candidate’s own future organic spending – a gain the wait branch picked up automatically (it is fit against simulated future states that already reflect the grown D), while the deploy branch’s analytical shortcut did not. Adding the trickle’s expected drift to the deploy branch’s convolution before evaluating it produced the +1.3-to-+1.7 figures an earlier draft of this section reported.

A variance double-count, found in an external review of this section, resolved in two passes. A reviewer questioned whether Section 4.4’s remaining-idiosyncratic-uncertainty proxy $V_i(t)$ and Paper I’s static σ_i were being combined correctly, or whether σ_i^2 was implicitly counted twice. Tracing the implementation confirmed the concern precisely: the simulator’s $\mu_i(t)$ already embeds the resolved-to-date share of the idiosyncratic budget via $\varepsilon_i^{\text{cum}}(t)$, so the deploy-branch convolution’s $\sqrt{\sigma_i^2 + V_i(t)}$ term was pricing the same uncertainty a second time (Section 4.5’s addendum derives the mechanism and verifies the telescoping identity numerically). Correcting this at every *intermediate* period – using $\sqrt{V_i(t)}$ alone, since $\mu_i(t)$ already reflects what has resolved – moved $\Theta(0)$ from the +1.3-to-+1.7 range to a +4.6-to-+5.9 range. The correction was verified three ways before being trusted at that stage: (1) a direct numerical check that $\text{Var}(\varepsilon_i^{\text{cum}}(t)) + V_i(t)$ is exactly constant across t (confirming the telescoping identity the fix depends on); (2) the pre-existing simulator self-consistency gate (Table 8) re-run and unchanged, since that gate validates the forward simulator’s calibration rather than the deploy-branch convolution where the error lived; and (3) a purpose-built mechanism-decomposition check (Section 8.7) in which a “nothing evolves” benchmark scenario, which should trivially give $\Theta(0) = 0$, in fact returned $\Theta(0) = 0.000$ only after a related boundary-condition inconsistency in that decomposition’s own implementation was also found and fixed.

A second round of external review of this same section then caught that the fix above had not been applied at the $t = T$ terminal boundary itself for the default (stochastic-shocks-on) code path – the anchor value seeding the entire backward induction still used σ_i rather than the same $V_i(t) \rightarrow 0$ limit every other period now used, an inconsistency between the corrected main text (which already stated the boundary should be handled by a limit, not by reintroducing σ_i) and what the code that produced Table 11’s figures actually did. Making this uniform – the terminal condition now uses the identical formula as every other period, which is exactly 0 variance at T by construction, requiring no special case at all – moved $\Theta(0)$ to the +4.5-to-+5.1 range reported in Table 11, a modest further change (a few tenths of a seat per scenario) rather than another multiple-fold swing. This second pass is reported with the same discipline as the first: the mechanism-decomposition sanity check (Section 8.7) was re-run and continued to confirm $\Theta(0) = 0.000$ for the “nothing evolves” benchmark, and every downstream table in this paper (11 through 13b, and Section 8.8) reflects this fully corrected version.

Corrections are reported in full, including their magnitude and direction, rather than silently folded into a single “final” number, because the discipline of checking a surprising result’s mechanism before publishing it is itself part of this paper’s contribution – the same discipline that, in Paper I, caught a materially different implementation bug in the marginal-seat-gain gradient.

8.7 Mechanism Decomposition: Isolating Information Value from Deterministic Sequencing

A reviewer raised a further, distinct concern: that $\Theta(0)$ ’s positive sign might reflect predictable candidate-spending growth (a deterministic sequencing/crowd-out effect) rather than genuine information-option value, since the sign flip documented above occurred specifically when the trickle channel was added. This is answered directly by re-solving $\Theta(0)$ five times, at the live 98-day horizon under the `eta_bootstrap_all_cycles` calibration, independently toggling each of the three stochastic-transition-law components off:

Table 13b: Mechanism Decomposition of $\Theta_{LP}(0)$ (LP Control Approximation – Suspended, Section 8.9)

Scenario	Trickle	Stochastic shocks	Opponent reaction	$\Theta(0)$
A: static benchmark	Off	Off	Off	+0.000
B: deterministic sequencing alone	On	Off	On	+3.479
C: pure information alone	Off	On	Off	+1.485
D: information + growth	On	On	Off	+5.071
E: full model	On	On	On	+4.714

Scenario A – nothing evolves over the horizon at all – gives exactly $\Theta(0) = 0$, the sanity check the decomposition machinery must pass before the other four rows can be trusted; it still does after the terminal-boundary correction described above. Scenarios B and C isolate each channel on its own: **deterministic sequencing value (B, +3.479) is larger than pure information value (C, +1.485)**, confirming the reviewer’s suspicion directly rather than refuting it. Roughly 70% of the two channels’ combined single-channel value ($3.479/(3.479+1.485)$) comes from candidate committees’ predictable organic spending growth, not from resolving genuine uncertainty – essentially unchanged from before the terminal-boundary correction, because Scenarios A and B both run with `enable_stochastic=False` and were therefore unaffected by it; only C, D, and E (which use the stochastic path where the boundary inconsistency lived) moved, and none by enough to change this conclusion. Scenario D (both channels, no opponent reaction) exceeds the sum of B and C, indicating the two channels interact super-additively rather than simply adding; adding opponent reaction (E) then reduces $\Theta(0)$ from D’s +5.071 to the full model’s +4.714, since reaction dampens the wait branch’s advantage by having $R_{i,t}$ partially offset the trickle it drives.

The honest claim, and the one this decomposition does not license. Per the reviewer’s own proposed resolution: the defensible claim is that this model estimates a positive value of deferring party deployment under *both* predictable candidate-spending growth *and* resolving uncertainty, with the deterministic channel contributing the larger share in isolation. The stronger claim – that Θ specifically prices information-driven flexibility – is not supported by this decomposition; Scenario C shows genuine information value is real and positive (+1.485) but is not, on its own, the dominant driver of the headline result.

8.8 Statistical Rigor: Simulation Noise and Out-of-Sample Policy Evaluation

The headline $\Theta(0)$ figures above are single-seed point estimates with no simulation-noise accounting, and the continuation-value regression is fit and evaluated on the same $K = 2,000$ paths – the standard in-sample look-ahead-bias risk in Longstaff–Schwartz applications. Both are addressed directly, on the `eta_bootstrap_all_cycles` calibration:

Monte Carlo standard error. Re-solving $\Theta(0)$ at $K = 2,000$ across 5 independent seeds gives a mean of +4.740 with sample SD 0.099 and standard error 0.044 expected seats – simulation noise is under 1% of the point estimate, not a material source of uncertainty relative to the roughly 3–4× swing Section 8.6’s variance-double-count correction produced (the terminal-boundary refinement that followed it moved figures by a few tenths of a seat, a magnitude these same checks below also cover).

Out-of-sample policy evaluation. Refitting with 30% of paths held out of every period’s continuation-value regression (so a held-out path’s own realized future value never informs the regression that decides its stopping choice) gives $\Theta(0) = +4.599$ on the held-out paths versus +4.604 in-sample – a 0.1% difference. The compressed, five-feature regression basis (Section 4.6) does not appear to be meaningfully overfitting at this sample size.

K-sensitivity. $\Theta(0) = +4.644$ at $K = 2,000$ versus +4.769 at $K = 5,000$ (single seed) – consistent with the multi-seed SE above, and well within the range needed to trust the headline figures at the

precision reported.

None of these checks change the qualitative or quantitative conclusion; they establish that remaining Monte Carlo and regression-specification uncertainty is small relative to $\Theta(0)$'s magnitude, which is not the same claim as small relative to the *calibration and parameter* uncertainty discussed in Limitations below.




Figure 8: Monte Carlo and regression-specification diagnostics, `eta_bootstrap_all_cycles`: (A) five-seed agreement, (B) K-sensitivity, (C) held-out policy evaluation against in-sample.

Section 8.8's checks establish that Monte Carlo noise and regression specification are small sources of uncertainty *for the LP-based calculation specifically*. They say nothing about whether the LP allocator itself is the right thing to be running this uncertainty analysis on. Section 8.9 shows it is not.

8.9 A Critical, Unresolved Allocator-Robustness Finding

Every Θ figure reported so far – Tables 11 through 13b, and Section 8.8's diagnostics – uses the fast LP allocator (`optimize()`) for the deploy branch's within-period spending decision, a deliberate substitution for Monte Carlo tractability (Appendix E.1). Appendix E.1 already documents that this substitution is not innocuous: the LP treats each race's marginal seat gain as a fixed constant, has no mechanism for diminishing returns, and degenerates into a greedy knapsack that concentrates money into a handful of races up to a cap rather than spreading it the way the true, concave objective would. That was reported as a specification difference requiring further validation, not yet as a threat to the paper's headline conclusion.

It is a threat to the paper's headline conclusion. A reduced-scope robustness check replaces the LP allocator with the true nonlinear allocator (`optimize_nonlinear()`, which correctly respects the persuasion ceiling and diminishing returns) throughout the *entire* backward induction – not only at the immediate $t = 0$ decision, but at every period the wait branch might eventually deploy – and re-solves $\Theta(0)$ for the `eta_bootstrap_all_cycles` scenario. Because `optimize_nonlinear()` costs roughly 40 seconds to over an hour per call (versus milliseconds for the LP), this check uses a drastically reduced path count ($K = 15$ versus the headline $K = 2,000$), with identical simulated

opponent-response draws used for both allocators so the comparison isolates the allocator's own effect:

Table 13c: LP-Throughout vs. Nonlinear-Throughout, Same Simulated Draws, $K = 15$ (First Pass)

	$\Theta(0)$ (expected seats)	frac_deploy_now(0)
LP-throughout	+5.097	0.0%
Nonlinear-throughout	−0.376	66.7%

The LP-throughout figure at $K = 15$ (+5.097) closely matches the headline $K = 2,000$ figure for this scenario (+4.692), confirming the reduced path count alone is not driving the difference below. The nonlinear-throughout figure is negative, with the stopping decision genuinely split across simulated paths (two of three favor deploying now) rather than the LP-throughout run's unanimous hold.

A second pass with common random numbers confirms this is the allocator, not simulation-path noise. Table 13c's two rows shared the same $\hat{\eta}/\text{resid}$ draws but not the same idiosyncratic-shock, R-reaction-noise, and national-environment state paths – a genuine gap identified in item (1) of the investigation plan below and closed immediately rather than left open. Resetting the simulator's random-number generator to an identical seed before each of the two `run_lsm()` calls (verified, before trusting any number, by confirming two LP-only runs under this reset scheme return bit-identical results) makes both allocators evaluate the exact same simulated $d_{i,t}$, $r_{i,t}$, $\mu_{i,t}$, G_t , and $\varepsilon_i^{\text{cum}}(t)$ paths, isolating $\Delta_{\text{allocator}} = \Theta_{\text{nonlinear}}(0) - \Theta_{\text{LP}}(0)$ from everything else that differs between the two runs:

Table 13d: LP-Throughout vs. Nonlinear-Throughout, Common Random Numbers, $K = 15$ (Paired)

	$\Theta(0)$ (expected seats)	frac_deploy_now(0)
LP-throughout	+4.817	0.0%
Nonlinear-throughout	−0.347	66.7%
$\Delta_{\text{allocator}}$	−5.164	—

Both figures land within a few tenths of a seat of the unpaired first pass (+5.097 $\rightarrow$ +4.817; −0.376 $\rightarrow$ −0.347), and the paired gap (−5.164) is essentially the same magnitude as the unpaired gap (−5.473). This is informative on its own: if the reversal were an artifact of the two runs happening to simulate different worlds, pairing should have moved the gap substantially; instead it barely moved at all, which is evidence the effect really is the allocator choice rather than which random states happened to get drawn. **A preliminary allocator-robustness analysis, now confirmed under common random numbers, indicates that the estimated value of waiting is highly sensitive to the within-period control solver. At $K = 15$, paired on identical simulated state paths, the LP-based backward induction produces $\Theta(0) = +4.817$, while a nonlinear-throughout implementation produces $\Theta(0) = -0.347$**

($\Delta_{\text{allocator}} = -5.164$). The nonlinear estimate is still too imprecise at this path count to establish the sign of $\Theta(0)$ with confidence, but the magnitude and now-demonstrated stability of the discrepancy make clear that the LP-based estimate cannot be interpreted as a policy-robust value of waiting. The paper therefore suspends its prior hold recommendation pending higher-powered nonlinear evaluation.

A third pass adds two more independent seeds – the allocator effect is now confirmed stable across simulated worlds, not just one. Item (2) of the investigation plan calls for checking $K = 15$ across multiple seeds before paying for a larger K . Two additional independent replicates (each with its own $\hat{\eta}$ /resid bootstrap draw and its own state-path RNG seed, LP and nonlinear paired within each replicate exactly as Table 13d’s methodology) give:

Table 13e: $\Delta_{\text{allocator}}$ Across Three Independent Seeds, $K = 15$ (Paired Within Each Seed)

Seed	$\Theta_{\text{LP}}(0)$	$\Theta_{\text{nonlinear}}(0)$	$\Delta_{\text{allocator}}$
20260730	+4.817	−0.347	−5.164
1	+4.608	−0.115	−4.723
2	+3.963	−0.315	−4.278
Mean (SD)	+4.463 (0.362)	−0.259 (0.109)	−4.722 (0.443)

Two findings here matter more than the point estimates. First, $\Delta_{\text{allocator}}$ itself is tightly clustered (−4.28 to −5.16, SD 0.44) despite each seed simulating an entirely different random world – the allocator’s own effect is precisely estimated even at $K = 15$, which is not true of $\Theta_{\text{nonlinear}}(0)$ on its own. Second, and more consequentially: **all three seeds give a negative $\Theta_{\text{nonlinear}}(0)$** . A single negative estimate is consistent with noise around zero; three independent negative estimates, with a sample mean (−0.259) more than two sample standard deviations from zero, is harder to dismiss that way, though $n = 3$ is still too few seeds for a formal confidence interval and this remains short of item (2)’s full incremental- K program.

A period-by-period decomposition locates a concrete mechanism (item (4)). Comparing the two allocators directly on a deterministic, trickle-only representative trajectory (no stochastic terms, so this needs only four cheap nonlinear calls rather than a full Monte Carlo run) at four points across the live horizon:

Table 13f: LP vs. Nonlinear Deploy Value and Allocation Pattern, Deterministic Trajectory

Days remaining	LP deploy value	Nonlinear deploy value	Diff	LP: races funded / top-5 share	Nonlinear: races funded / top-5 share
98	230.53	233.74	+3.20	7 / 75.0%	237 / 12.3%
70	232.25	235.22	+2.97	7 / 75.0%	244 / 11.5%
42	233.05	235.15	+2.10	7 / 75.0%	241 / 11.0%

Days remaining	LP deploy value	Nonlinear deploy value	Diff	LP: races funded / top-5 share	Nonlinear: races funded / top-5 share
14	231.92	233.01	+1.09	7 / 75.0%	258 / 10.3%

The LP allocator funds exactly 7 races at every single point on the horizon, always concentrating 75% of the party budget in its top 5 – a static, state-insensitive pattern, not a dynamically adjusting one. The nonlinear allocator funds 237–258 races throughout, with concentration falling slightly as Election Day approaches (12.3% $\rightarrow$ 10.3%). More importantly, the *gap* between the two allocators’ deploy value shrinks steadily and substantially as the horizon shortens: +3.20 seats at 98 days out, down to +1.09 at 14 days out – a roughly 3 $\times$ narrowing. This gives the mechanism a specific, falsifiable shape rather than a general plausibility argument: deploying immediately is penalized *most* by the LP approximation far from Election Day, and *least* penalized close to it, which is exactly the asymmetry that would make “wait, then deploy later” look artificially more attractive than “deploy now” under the LP – the deferred deploy option is being scored against a version of the allocator that is closer to its true, nonlinear value by the time it actually gets used.

$K = 30$ **confirms the effect does not wash out as precision increases.** The remaining half of item (2) is an actual increase in K , not more replicates at $K = 15$ – more seeds at a fixed, small K cannot substitute for reducing each individual estimate’s own noise. A single $K = 30$ replicate (double the path count, hence roughly double the compute cost, an independent seed and independent state paths from all $K = 15$ runs):

Table 13g: $K = 30$, Paired

K	$\Theta_{\text{LP}}(0)$	$\Theta_{\text{nonlinear}}(0)$	$\Delta_{\text{allocator}}$
15 (mean of 3 seeds)	+4.463	−0.259	−4.722
30 (1 seed)	+4.268	−0.435	−4.703

Doubling K moved $\Delta_{\text{allocator}}$ by less than half a percent (−4.722 $\rightarrow$ −4.703), and $\Theta_{\text{nonlinear}}(0)$ remained negative and within the range the three $K = 15$ seeds already spanned. This is the single most reassuring result in the investigation so far: if the reversal were a small- K artifact that would resolve toward zero or flip positive with more precision, doubling K should have moved it visibly. It did not.

What this result establishes. The estimated value of waiting is highly dependent on the within-period allocator, by an amount (roughly 4.3–5.5 seats, consistent across the unpaired comparison, the paired comparison, three $K = 15$ seeds, and now one $K = 30$ replicate) far too large to be ordinary Monte Carlo variation around the LP result, and stable under a real increase in path count. The LP-based calculation cannot, at present, support statements such as “holding is the unanimous recommendation,” “waiting is worth +4.6 to +5.9 expected seats,” “the result is robust across calibration scenarios,” or “the optimal policy favors holding the entire reserve.” Those are

results of the LP approximation, not yet demonstrated properties of the underlying nonlinear model this paper intends to specify. The period-by-period decomposition further establishes a specific, quantified mechanism: the LP allocator’s static, state-insensitive concentration in a handful of races penalizes immediate deployment more heavily earlier in the horizon than later, which by itself would inflate the apparent value of waiting.

What this result does not establish. No single replicate at $K = 15$ or $K = 30$ (Tables 13c–13g) individually establishes that immediate deployment is preferable, and four total replicates (three at $K = 15$, one at $K = 30$) is not enough to replace a proper, larger- K progression with multiple seeds at each level. The continuation regressions remain data-starved at these path counts (five portfolio-level features fit against a tiny effective sample); the stopping rule can change from a handful of paths; nonlinear-solver variation (SLSQP convergence differs across states, sometimes needing far more iterations than others) may add noise the LP runs do not have. A rough scaling of Section 8.8’s $K = 2,000$ LP standard error (0.044) by $\sqrt{2,000/K}$ suggests an approximate SE near 0.5 at $K = 15$ and 0.35 at $K = 30$ – broadly consistent with the observed cross-seed SD of 0.109 on $\Theta_{\text{nonlinear}}(0)$ at $K = 15$, though this scaling assumes the same estimator with stable asymptotic behavior, which the nonlinear control’s different regression targets and stopping decisions may not fully satisfy. Pairing, the multi-seed check, and the $K = 30$ replicate each addressed a different specific source of doubt (simulated-world artifact; lucky/unlucky single draw; small- K instability that should wash out with more precision) and none of them moved the answer; none of this yet substitutes for out-of-sample policy evaluation or multiple seeds at $K = 30$ and beyond. The defensible conclusion at this stage is that the nonlinear estimate is consistent with a small negative or near-zero value of waiting, increasingly rather than merely “not positive” – but still short of a confident sign determination.

Why the LP could be manufacturing artificial option value. Table 13f’s period-by-period decomposition turns what was originally a plausibility argument into a specific, quantified mechanism. The LP allocator’s greedy, corner-concentrated behavior distorts the deploy branch: it funds exactly 7 races to their caps at every point across the horizon, 75% of the budget concentrated in the top 5, regardless of how much time remains – a static pattern showing no state-sensitivity at all. Because this static concentration is a worse approximation to the true, concave objective far from Election Day than close to it (Table 13f’s deploy-value gap narrows from +3.20 seats at 98 days to +1.09 at 14 days), immediate deployment is penalized more heavily the earlier it is evaluated – exactly the asymmetry that would make “wait, then deploy once less time remains” look artificially more attractive than “deploy now,” for a reason having nothing to do with genuine information value or organic spending growth. In other words, the LP appears to manufacture apparent optionality through its own greedy, state-insensitive allocation pattern rather than through anything a rational, correctly-optimizing committee would actually experience.

What remains to be done before this is resolved. A structured investigation is required before $\Theta(0)$ ’s sign under the intended nonlinear model can be reported with confidence. Item (1), paired simulation using common random numbers, **is done** (Table 13d): the reversal survives pairing. Item (2)’s first sub-step, multiple seeds at $K = 15$, **is done** (Table 13e): three independent seeds all show the same reversal, tightly clustered. Item (4), a period-by-period mechanism decomposition,

is done (Table 13f): the mechanism is the LP’s state-insensitive concentration penalizing early deployment more than late. Item (2)’s next step, an actual increase in K , **has begun** (Table 13g): one $K = 30$ replicate shows the same effect at essentially the same magnitude. Remaining: multiple seeds at $K = 30$ (matching the rigor already applied at $K = 15$, not just the single replicate so far), then $K = 50$ and $K = 100$ if $K = 30$ ’s pattern holds; (3) out-of-sample policy evaluation (Section 8.8’s held-out-path methodology) applied to the nonlinear-throughout run once K is large enough to support it; (5) a validated fast surrogate that preserves diminishing returns (a piecewise-linear concave approximation, a precomputed per-race value-on-a-spending-grid table, or a model trained on nonlinear-optimizer outputs) validated against the true nonlinear solver before being trusted for a full $K = 2,000$ -scale re-solve; and (6) only after the `eta_bootstrap_all_cycles` scenario stabilizes, repeating across `eta_fit_2022`, `eta_fit_2024`, and the continuous deployment-fraction framing – the 364-day comparative-static extension is lower priority still. This section now reports a substantially more complete case that the reversal is real and stable under increasing precision, plus its mechanism – but not yet a large-enough- K resolution of $\Theta(0)$ ’s sign under the intended nonlinear model, and the paper’s hold recommendation remains suspended.

9 Discussion

9.1 Why the Corner Flipped (Within the LP Approximation – Itself Now Suspended)

This subsection describes the LP-based calculation’s own internal history and should be read alongside Section 8.9’s finding that the LP figures it explains are not currently trustworthy as a policy conclusion. The mechanism is structural, not a recalibration accident. Once candidate-committee spending is allowed to grow deterministically while a committee waits, and opponent reaction is correctly credited with responding to that growth, the wait branch’s simulated future states genuinely differ from – and, on average, exceed – what the deploy branch’s convolution identity captures for a static $D_{i,t}$. The earlier, “deploy favored” result was not wrong because the option-value logic of Papers II–III is wrong; it was wrong because the model actually tested was missing a real channel through which waiting pays off. This is a direct illustration of Section 3.3’s methodological claim: a Bellman equation over an incompletely specified P does not merely give an imprecise answer, it can give the *wrong-signed* one. The subsequent variance-double-count correction (Sections 4.5, 8.6) is a second illustration of the same claim, one level down: even with the transition law’s *components* correctly specified, an error in how two of their variance terms combine changed $\Theta_{LP}(0)$ ’s *magnitude* by a factor of roughly 3–4, without changing its sign. Section 8.9’s allocator-robustness finding is a third, more consequential illustration still: a specification error in the *control* (the within-period allocator), not the transition law P this paper is nominally about, appears large enough to change $\Theta(0)$ ’s sign a second time, and this one is not yet resolved.

9.2 What Remains Structurally Unpriced – and What the Decomposition Shows Is Already Priced

Even the corrected model has a real, stated scope boundary. The convolution identity (Section 4.5) is exact for what it integrates over – the expected effect of unresolved shocks on terminal win probability – but it does not, and structurally cannot, capture the value of *adaptive* decision-making: that new information arriving mid-campaign could change *how* the reserve is allocated, not merely *whether* it is deployed. Deployment in this model is a one-time absorbing decision, computed once, at the moment of deployment. $\Theta(t)$ throughout this paper should therefore be read not as “the value of flexibility” in a fully general sense, but as **the value of flexibility net of what the deploy branch’s own analytical widening already captures for free** – a narrower and more accurate claim.

Section 8.7’s mechanism decomposition sharpens this further in a different direction: even within the narrower object Θ actually measures, a majority of its magnitude in isolation (Scenario B versus Scenario C) is deterministic-sequencing value – the predictable fact that candidate committees keep spending their own money while the party waits, which a myopic “deploy now and never revisit” benchmark does not credit itself with – rather than information-option value in the classical real-options sense. This does not make Θ a less real or less useful quantity for the decision at hand; a committee choosing whether to hold or deploy today should care about the total value of waiting, whatever its source. But it does mean the paper’s own framing needs a qualifier: Θ prices *the*

value of deferring commitment given both predictable organic growth and resolving uncertainty, not information value on its own, and the two should not be conflated when this result is cited elsewhere.

9.3 Reconciling with Paper II’s Baseline

Paper II’s own live run, using the Θ -free rollout architecture, recommends deploying the full reserve immediately, concentrated disproportionately in non-competitive seats. The LP-based version of this paper’s result recommended the opposite action at the same decision point, and the two were framed as exactly the comparison this research program is built to make: Paper II identifies and measures the size of a gap; this paper specifies the process that determines which side of that gap the live decision actually falls on. Section 8.9 complicates this reconciliation directly – with the nonlinear allocator, this paper’s own preliminary estimate moves back toward Paper II’s “deploy” recommendation, at least at the reduced precision so far available. Which side of the gap the live decision actually falls on is, at present, unresolved by either paper, not settled in Paper III’s favor.

9.4 A Domain-Intuition Check Against Real Data

A model recommendation this consequential should not be accepted on the strength of internal consistency alone. Real independent-expenditure committees spent only 1.6–4.4% of their eventual full-cycle total by September 1, in both the 2022 and 2024 cycles and both parties – near-total real-world back-loading, and strong independent evidence that sophisticated political spenders behave nothing like an immediate-deployment strategy, for reasons this model may or may not fully capture (genuine information value, strategic waiting, targeting immaturity early in a cycle, or some mix). The LP-based result reported in Section 8.3–8.5 is consistent with this pattern; the earlier, uncorrected LP-based result directly contradicted it. This corroborating evidence should be read with the same caution Section 8.9 now attaches to the LP-based figures themselves: real committees’ backloading is independent evidence worth weighing, but it cannot substitute for resolving which allocator’s $\Theta(0)$ estimate is actually correct, and it is not, on its own, a reason to prefer the LP-based conclusion over the nonlinear one.

9.5 Generalizability

The specification-before-optimization discipline this paper insists on – that an optimal-stopping problem is not well-posed until its underlying stochastic process is specified, and that specifying the process is the actual scientific content of the problem – generalizes to any domain where a real-options framing is invoked without a calibrated transition law behind it. The three-component decomposition (a reactive/strategic process, an exogenous environmental process, and an idiosyncratic decay process) is likewise a template: any capital-commitment decision with a competitive response, a common macro factor, and unit-specific resolving uncertainty admits the same three-part specification, whether the underlying asset is a campaign, a venture investment, or a real physical option to expand capacity.

10 Limitations

10.1 Allocator Robustness (Critical, Unresolved – See Section 8.9)

Every limitation below describes the LP-based calculation. That calculation’s own headline conclusion is currently suspended: Section 8.9 finds that replacing the LP within-period allocator with the true nonlinear allocator throughout the backward induction reverses $\Theta(0)$ ’s sign at a reduced path count, by a margin (roughly 5.5 seats) far exceeding this paper’s own Monte Carlo noise estimate. This is listed first, not folded into “Computational” below, because it is not a precision issue – it is an open question about which sign is correct, and it takes priority over every other limitation in this section until resolved.

10.2 Data

Race-level idiosyncratic uncertainty ($\varepsilon_{i,t}$) is a bounded proxy, not a fitted process, for a data reason this paper judges likely permanent under a public-data-only constraint: district-level polling density is too sparse, for the large majority of House races, to support genuine time-series estimation. Opponent reaction is estimated from IE-to-IE reaction and then applied, in the deploy-branch drift correction, to a candidate-committee spending channel it was not originally estimated on – an untested extension of an existing estimate, flagged explicitly rather than assumed to transfer. Five of seven opponent-reaction tiers show statistically significant cycle-to-cycle variation rather than a single stable constant, so the cycle-weighted point estimates in Table 5 should be read as averages over real historical variation, not as fixed structural parameters.

10.3 Modeling

The random-walk versus mean-reversion choice for G_t is empirically underdetermined at the three-to-nine-month horizon that matters most for this paper’s live application; a longer-horizon application would need to resolve this distinction more carefully. The realized late-cycle generic-ballot asymmetry (Section 8.1’s Validation B) is directionally consistent across four cycles but not reliable at $n = 4$, and is not fed into μ_i ’s structural formula regardless, per the scope boundary of Section 4.7. The deploy-branch convolution identity, as discussed in Section 9.2, prices a narrower object than the full value of adaptive flexibility, by construction.

10.4 Computational

The Longstaff–Schwartz regression basis is compressed to five portfolio-level features specifically to avoid overfitting at the sample sizes involved (Section 4.6); this is standard practice for this class of problem but is a genuine information loss relative to a (computationally infeasible) full per-race basis. Section 8.8 directly quantifies remaining Monte Carlo and regression-specification uncertainty: a 5-seed standard error of 0.044 expected seats (under 1% of the point estimate), a 0.1% in-sample-vs.-held-out gap under a 30%-held-out policy evaluation, and consistency between $K = 2,000$ and $K = 5,000$. Monte Carlo noise remains visible at the margins of the continuous-framing results (Table 13’s `eta_bootstrap_all_cycles` scenario, where non-hold mass is now spread across several interior grid fractions rather than concentrated at one); an 11-point grid spot-check of the cor-

rected continuous- ϕ result on both single-cycle brackets (`eta_fit_2022`, `eta_fit_2024`) confirms the 5-point figures directly (differences of 0.008 and 0.011 seats respectively, Section 8.5), though the `eta_bootstrap_all_cycles` scenario’s own interior-mass distribution – the one genuinely non-corner result in this research line – has not itself been re-checked at 11 points and remains the more precision-sensitive of the three.

This Monte Carlo standard error is not the paper’s principal source of uncertainty around $\Theta(0)$, and should not be read as one. It quantifies simulation noise alone, holding every calibrated parameter fixed at its point estimate. The larger, unquantified uncertainty comes from those point estimates themselves: the candidate-spending trickle rate (Section 6.4), the opponent-reaction coefficients $\hat{\eta}(\text{tier})$ (Table 5, several tiers of which show real cycle-to-cycle instability rather than a stable constant), the decay rate λ borrowed from the national-environment term structure and applied to the idiosyncratic-uncertainty proxy (Section 4.4), and Paper I’s underlying margin-model coefficients. None of these are re-estimated with bootstrap resampling and propagated through to $\Theta(0)$; doing so would require a nested parameter bootstrap – refitting each calibration input on resampled historical panels and re-solving $\Theta(0)$ under each draw – which this paper does not attempt. The 0.044-seat simulation-noise figure above answers “how much does re-running the same calibrated model change the answer,” not “how much would a different, equally defensible calibration change the answer,” and the two should not be conflated when this paper’s precision is cited elsewhere.

10.5 Practical Deployment

No practical recommendation is currently made. An earlier version of this section stated that the LP-based finding – hold the reserve at the live 2026 decision – was a recommendation about a single, specific decision point, re-derived from the current state each time the pipeline is re-run. Section 8.9 supersedes that framing: until the allocator-robustness question is resolved, this paper does not recommend either holding or deploying the live 2026 reserve. Paper II’s own architecture (Section 4.7 of that paper) already anticipates that a Θ -aware recommendation must be recomputed as the state evolves, not applied once and left unrevisited; that design is unaffected by the present suspension, but there is presently no Θ -aware recommendation to feed into it. This paper’s Longstaff–Schwartz solve has not been folded back into Paper II’s rollout loop as a per-period, automatically-applied reserve fraction, and should not be until Section 8.9’s investigation concludes.

11 Conclusion

Paper I answers what one more campaign dollar is worth today. Paper II asks how to make that decision repeatedly, and discovers it cannot yet answer whether to wait. This paper does not answer that question directly; it answers the prior question that determines whether it is even askable: how does the political state evolve on its own, absent any decision. Θ , reserve policy, and optimal stopping are all corollaries of that answer, in exactly the sense that an option’s value is a corollary of a specified stock-price process rather than an independently derived quantity. The scientific contribution of this research line is not “we solved a Bellman equation.” It is “we estimated a defensible transition law P for the political state” – and, honestly, in one component (race-level idiosyncratic uncertainty), “we established that this cannot currently be estimated from public data, and said so plainly.”

Having specified and estimated P , solved the resulting Bellman equation, and corrected real errors found by checking mechanism before trusting output – a missing spending-drift channel, an asymmetric convolution, and a variance double-count found in external review and resolved in two passes as that review’s own scrutiny deepened – this paper’s LP-based calculation found $\Theta_{LP}(0)$ substantially positive (+4.5 to +5.1 expected seats) in every calibration scenario tested, in both a binary and a continuous deployment-fraction framing, at both a 98-day horizon and a 364-day comparative-static extension, with holding the deployable reserve dominating deploying it. **That finding is no longer this paper’s conclusion.** Section 8.9 reports a further discovery, made by turning the same mechanism-checking discipline on the control itself rather than only on the transition law P : replacing the LP within-period allocator with the true, diminishing-returns-respecting nonlinear allocator throughout the backward induction reverses $\Theta(0)$ ’s sign at a reduced path count, by a margin far exceeding this paper’s own quantified Monte Carlo noise. A mechanism decomposition (Section 8.7) had already shown that a majority of the LP-based Θ ’s magnitude, isolated channel by channel, traces to predictable candidate-spending growth rather than classical information-option value – itself a real, quantified qualification of the paper’s original framing – but Section 8.9’s finding is more fundamental: it is not yet known which sign the correctly specified nonlinear model implies, and this paper does not currently endorse either “hold” or “deploy” as its live 2026 recommendation.

This is not a failure of the paper’s central methodological claim; it is a further instance of it. Section 3.3 argues that an optimal-stopping calculation is not well-posed until its underlying transition law is specified – that specifying P is the actual scientific content of the problem, not the machinery that follows from it. Section 8.9 shows the same discipline extends to the *control*: an optimal-stopping calculation is equally not trustworthy until the within-period optimization it depends on is verified to be the one actually intended, not a tractable substitute assumed innocuous. This paper specified and estimated P successfully. It has not yet verified that its reported Θ is a property of the intended nonlinear control rather than an artifact of the LP control’s own greedy, discontinuous behavior. Together, the three papers in this research program price the next dollar (Paper I), build the system that deploys it repeatedly (Paper II), and attempt to price the alternative of not yet deploying it (this paper) – and, at the specific decision this program was built to inform, this paper currently has no defensible answer to give, pending the investigation Section 8.9 lays out.

12 Data Availability

All data used in this paper are drawn from public sources: the sources already documented in Papers I and II, plus the FEC API’s dated committee-reports endpoint (`/committee/{id}/reports/`), the FEC comprehensive independent-expenditure bulk export, and a Wayback Machine snapshot of FiveThirtyEight’s discontinued generic-ballot data feed, whose provenance is documented in the project’s data catalog.

13 Code Availability

Repository: <https://github.com/callum-doty/political-portfolio> **Entry points:** `scripts/estimate_eta_reaction.py`, `scripts/estimate_gb_volatility.py`, `scripts/estimate_candidate_` (calibration); `scripts/solve_bellman_lsm.py` (binary framing, including the mechanism-decomposition toggles and out-of-sample `held_out_frac` option used in Sections 8.7–8.8); `scripts/solve_bellman_lsm_continuous_phi.py` (continuous framing); `scripts/theta_mechanism_decomposi` (Section 8.7); `scripts/theta_statistical_rigor.py` (Section 8.8); `scripts/validate_state_simulator.py` (Section 8.1’s validation suite); `scripts/simulate_and_validate.py` (Table 8’s self-consistency checks); `tests/test_bellman_lsm.py` (automated regression tests covering the corrections of Section 8.6)

14 Conflict of Interest

The authors declare no conflict of interest. This research was not funded by, and the authors hold no financial relationship with, any political campaign, party committee, or campaign consulting firm.

15 References

- Andersen, T. G., and Bollerslev, T. (1998). Answering the skeptics: Yes, standard volatility models do provide accurate forecasts. *International Economic Review*, 39(4), 885–905.
- Black, F., and Scholes, M. (1973). The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3), 637–654.
- Dixit, A. K., and Pindyck, R. S. (1994). *Investment Under Uncertainty*. Princeton University Press.
- Erikson, R. S., and Palfrey, T. R. (2000). Equilibrium in campaign spending games. *American Political Science Review*, 94(3), 595–609.
- Gelman, A., and King, G. (1993). Why are American Presidential election campaign polls so variable when votes are so predictable? *British Journal of Political Science*, 23(4), 409–451.
- Kalman, R. E. (1960). A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1), 35–45.
- Longstaff, F. A., and Schwartz, E. S. (2001). Valuing American options by simulation: A simple least-squares approach. *The Review of Financial Studies*, 14(1), 113–147.
- McDonald, R., and Siegel, D. (1986). The value of waiting to invest. *The Quarterly Journal of Economics*, 101(4), 707–727.
- Merton, R. C. (1973). Theory of rational option pricing. *The Bell Journal of Economics and Management Science*, 4(1), 141–183.
- Uhlenbeck, G. E., and Ornstein, L. S. (1930). On the theory of the Brownian motion. *Physical Review*, 36(5), 823–841.
- Machine-readable BibTeX entries are provided as `references.bib` in the replication repository, combined with Paper I’s and Paper II’s reference lists.

16 Appendix A: Notation Reference Table

See Table 2 (Section 3.1) for symbols specific to this paper. Every symbol inherited from Paper I (the spending response surface, μ_i , σ_i , MSG_i , Φ , φ) and Paper II ($\mathbf{X}_t$ ’s originally-informal components, B_t , L_t , F_t , λ ’s original EMA usage) retains its prior definition unless explicitly redefined here.

17 Appendix B: Derivations

17.1 B.1 The Deploy-Branch Terminal-Probability Identity (Corrected)

An earlier version of this appendix derived the deploy branch’s win probability as a genuine Gaussian convolution: for $Z \sim N(0, 1)$ and constants a, b with $b > 0$, $\mathbb{E}_Z[\Phi(a + bZ)] = \Phi(a/\sqrt{1 + b^2})$, applied with $\mu_i(T) = \mu_i(t) + \xi$, $\xi \sim N(0, V_i(t))$, and $P(\text{win}_i) = \Phi(\mu_i(T)/\sigma_i)$ to give $\Phi(\mu_i(t)/\sqrt{\sigma_i^2 + V_i(t)})$. That derivation treated $\mu_i(t)$ as a clean structural mean with no resolved shock embedded in it, and σ_i as a separate, permanent noise floor sitting on top of the future shock ξ . Section 4.5’s addendum and Section 8.6 report why this double-counts: the implemented simulator’s $\mu_i(t) = \mu_i^{\text{struct}}(t) + \varepsilon_i^{\text{cum}}(t)$ already embeds the resolved-to-date share of the idiosyncratic-uncertainty budget, and $\text{Var}(\varepsilon_i^{\text{cum}}(t)) + V_i(t)$ is constant in t (Appendix B.2) – so σ_i^2 is not a separate quantity left over to add a second time.

The correct derivation is simpler than a convolution, precisely because there is no second, independent noise source left to integrate over. If $\mu_i(T) = \mu_i(t) + \xi$ for the genuinely mean-zero remaining shock $\xi \sim N(0, V_i(t))$ – the only movement not yet realized as of t , given that $\mu_i(t)$ already reflects everything resolved up to that point – and a race is “won” exactly when the fully resolved margin $\mu_i(T)$ is positive, then

$$P(\text{win}_i \mid \text{info at } t) = P(\mu_i(t) + \xi > 0) = P(\xi > -\mu_i(t)) = \Phi\left(\frac{\mu_i(t)}{\sqrt{V_i(t)}}\right),$$

a direct Gaussian tail probability. This is the identity Section 4.5’s corrected formula and every backward induction in Section 7 use; Appendix C.1 derives its behavior at the $t = T$ boundary.

17.2 B.2 Incremental Decomposition of the Idiosyncratic-Uncertainty Schedule

Section 4.4 gives the *cumulative* remaining-uncertainty target, $V_i(t) = \sigma_i^2(1 - e^{-\lambda(T-t)})$. For a period grid $n = 0, \dots, N$ (period length Δ , N = Election Day), the per-step increment variance is $v_{i,n} = V_i(n) - V_i(n+1)$. Drawing $\varepsilon_{i,n+1} \sim N(0, v_{i,n})$ independently at each step reproduces the cumulative schedule exactly by telescoping: $\sum_{n=t}^{N-1} v_{i,n} = V_i(t) - V_i(N) = V_i(t)$, since $V_i(N) = 0$ (matching $\Theta(T) = 0$ at Election Day). This is an independent-increment process matched to a prescribed, shrinking variance schedule, not a Brownian bridge in the strict sense, since nothing conditions on a known terminal value.

17.3 B.3 EMA/OU Term-Structure Fit

Fitting $\text{Var}(\Delta G)(t) = A(1 - e^{-t/\tau})$ by nonlinear least squares to the empirical term structure of Table 6 gives $\hat{\tau} = 186.5$ days, $\hat{\lambda} = 1/\hat{\tau} = 0.00536 \text{ day}^{-1}$, used identically in both Section 4.3’s national-environment calibration and, by the borrowed-rate logic of Section 4.4, the idiosyncratic-uncertainty proxy – a single fitted decay rate serving both purposes, not two independently-tuned parameters.

18 Appendix C: Proofs

18.1 C.1 $\Theta(T) = 0$ (Boundary Condition)

At $t = T$ (Election Day), no further information can arrive before the outcome is realized, so $V_i(T) = 0$ for every race (Appendix B.2’s telescoping identity). Appendix B.1’s identity is taken as a limit here, not as a reintroduction of σ_i : as $V_i(t) \rightarrow 0^+$, $\Phi(\mu_i(t)/\sqrt{V_i(t)}) \rightarrow \mathbb{1}[\mu_i(t) > 0]$, the deterministic indicator of the sign of the fully resolved simulated margin – consistent with $\mu_i(T)$ already being that fully resolved value, since it embeds the complete idiosyncratic-uncertainty budget via $\varepsilon_i^{\text{cum}}(T)$ (Appendix B.2), leaving nothing left to be uncertain about at T . This is the same value the wait branch’s continuation regression must converge to, since there is no $t = T + 1$ to wait for: both the deploy branch’s terminal anchor and the wait branch’s continuation prediction are evaluated against the identical simulated $\mu_i(T)$ for every path. Hence $\Theta(T) = V_T^{\text{wait}} - V_T^{\text{deploy-now}} = 0$ by construction, the boundary condition Section 3.5 states and every backward induction in Section 7 is anchored to.

19 Appendix D: Additional Robustness Analyses

19.1 D.1 Leave-One-Cycle-Out Stability Check

Refitting $\hat{\eta}$ separately on 2022 only and on 2024 only, rather than pooled, shows the three tiers Section 6.1 flags as most reliable on precision grounds (Toss-Up, Lean D, Lean R) keep the same sign in both single-cycle fits, but magnitudes are not stable (Toss-Up: 0.730 fit on 2022 alone vs. 0.341 fit on 2024 alone) – consistent with the significant cycle-to-cycle variation the full seven-cycle joint test (Table 5’s I^2 column) later confirms with adequate statistical power. Likely R’s single-cycle estimate flips sign between the two fits (-0.017 vs. $+0.412$), and this tier’s pooled estimate should not be trusted as a stable constant.

19.2 D.2 Necessary-Condition Check Detail

Section 8.2’s Jaccard-overlap check compares the top-twenty model-recommended races at a 60-day-out and a 14-day-out historical snapshot, using the true nonlinear optimizer (not the LP path used for Monte Carlo tractability elsewhere in this paper) and holding the total DCCC party budget fixed and identical at both snapshots so that only the information available differs. Overlap of 0.54 (2022) and 0.67 (2024) indicates roughly one-third to one-half of the top twenty races change between the two snapshots – real, substantial turnover attributable to information arrival alone.

20 Appendix E: Sensitivity Analyses

20.1 E.1 LP-vs-Nonlinear Allocator Divergence in the Deploy Branch

The deploy branch’s within-period allocation uses Paper I’s fast LP solver (`optimize()`, milliseconds per call) rather than the full nonlinear solver (`optimize_nonlinear()`, roughly 19 seconds per call), a deliberate substitution for Monte Carlo tractability given $K \times N \times (\text{scenarios}) \approx 32,000$ total calls. Running both allocators on identical inputs confirms this is a real, understood specification difference, not an approximation error: the LP objective treats each race’s marginal seat gain as a fixed constant with no diminishing-returns mechanism, so it degenerates into a greedy knapsack that funds only a handful of the lowest-floor races to the per-race cap and funds every other race, across every tier, at zero – a materially different allocation pattern than the nonlinear solver’s, for reasons having nothing to do with the value of waiting. This was checked directly (not assumed) before the deploy-branch comparisons in Section 8 were reported, since silently attributing an optimizer-specification difference to the Θ calibration would have been a genuine confound.

20.2 E.2 Sensitivity of $\Theta(0)$ to the Codebase-Wide Corrections in Paper I

An independent audit of the shared codebase (Paper I’s persuasion-ceiling fix and an earlier σ_i -model correction) was checked for its effect on this paper’s headline live-horizon result, since the deploy branch’s margin computation shares code with Paper I’s estimation pipeline. The persuasion ceiling itself does not enter `margin_gradient()`’s computation directly; the only channel by which Paper I’s corrections reach this paper’s result is the LP-scaling fix to the deploy branch’s within-period allocator, which was confirmed not to move Table 11’s figures materially (live $\Theta(0)$: 1.328/1.719/1.619 expected seats, essentially unchanged from the pre-audit 7-cycle-trickle figures of 0.997/1.808/1.295 expected seats – directionally consistent, with the differences attributable to the σ_i -model correction’s uneven effect across scenarios’ underlying win-probability curves, not to the ceiling itself). These are the pre-2026-07-28-audit figures; Section 8.6 documents a subsequent variance-specification correction that supersedes them.

21 Appendix F: Hyperparameters

Table F.1: Key Configuration Parameters

Parameter	Value	Section
K (Monte Carlo paths)	2,000	7.1
Reporting cadence	14 days (biweekly)	5.3
Live horizon	98 days	8.3
Counterfactual horizon	364 days	8.4
Regression basis dimension	5	4.6
Continuous- ϕ grid	5-point $\{0, .25, .5, .75, 1\}$, spot-checked at 11-point	8.5
λ (decay rate)	0.00536 day^{-1}	6.2

Parameter	Value	Section
Random seed (headline figures)	20260716	7.4
Multi-seed MC error check	5 independent seeds, $K = 2,000$	8.8
K -sensitivity check	2,000 vs. 5,000	8.8

22 Appendix G: Pseudocode

See Algorithm 1 (Section 7.2) for the complete backward-induction procedure.

23 Appendix H: Database Schema

Table H.1: Key Processed Artifacts

File	Contents
data/processed/eta_uncertainty.json	Bootstrap and random-effects $\hat{\eta}$ distributions (Section 6.1)
data/processed/gb_dynamics.json	σ_G per- $\sqrt{\text{day}}$, λ , τ (single source of truth, Section 6.2)
outputs/theta_schedule.json	Live-horizon $\Theta(t)$ by scenario and period (Table 11)
outputs/theta_schedule_1yr_counterfactual.json	365-day counterfactual results (Table 12)
outputs/theta_schedule_continuous_phi_*.json	Continuous-framing results by scenario (Table 13)
outputs/simulator_validation_summary.json	Validations A–C (Tables 9–10)
outputs/simulator_self_consistency.json	Four-part self-consistency check (Table 8)
outputs/eta_seven_cycle_extension.csv	Per-cycle $\hat{\eta}$ fits, leave-one-out (Appendix D.1)
outputs/eta_seven_cycle_joint_test.csv	Cycle-variation joint significance tests (I^2 , Table 5)

24 Appendix I: Additional Figures

Figures 1–6 (Sections 6.1, 6.2, 6.3, 8.3, 8.1, 8.5) constitute the complete set of figures generated for this paper; no additional figures beyond those already presented in the main text are included here.

25 Appendix J: Additional Tables

Table J.1: Leave-One-Cycle-Out $\hat{\eta}$ Fits

Tier	Fit on 2022 only	Fit on 2024 only
Toss-Up	0.730	0.341
Lean R	0.387	0.183
Lean D	0.294	0.237
Likely D	−0.212	−0.101
Likely R	−0.017	0.412
Safe D	1.070	0.318
Safe R	−0.003	−0.165

26 Appendix K: Configuration Files

Relevant excerpt of `config.yaml`’s dynamic-allocation block (full file in the code repository, Code Availability):

```
dynamic:
  ema_lambda: 0.7
  reporting_cadence: "biweekly"
  period_days: 14
  election_day_2026: "2026-11-03"

theta:
  k_paths: 2000
  live_horizon_days: 98
  counterfactual_horizon_days: 364
  regression_basis_features: 5
  continuous_phi_grid: [0.0, 0.25, 0.5, 0.75, 1.0]
```

27 Appendix L: Reproducibility Checklist

- ☒ All data sources are public and cited, or documented as requiring a free registered API key (Section 5.1)
- ☒ Complete calibration and solver code is version-controlled and publicly available (Code Availability)
- ☒ Random seeds are fixed for all Monte Carlo procedures (seed 20260716; Section 8.8 additionally reports a 5-seed Monte Carlo standard error)
- ☒ The simulator passes a four-part self-consistency check before any Θ figure is trusted (Section 8.1, Table 8)
- ☒ The valuation chain underlying every simulated path is validated against real held-out outcomes (Section 8.1, Table 9)
- ☒ Three real implementation errors are reported in full, including their direction and magnitude, not silently corrected (Section 8.6: a missing spending-drift channel, an asymmetric convolution, and a variance double-count including its terminal-boundary consistency extension)

- ⊗ Continuation-value regression is evaluated out-of-sample on held-out paths, not only in-sample (Section 8.8)
- ⊗ The binary hold-or-deploy framing is tested against a genuinely continuous generalization, not assumed adequate (Section 8.5)
- ⊗ Results are re-verified against the live codebase’s independent corrections (Appendix E.2)
- **CRITICAL, UNRESOLVED:** the LP-vs-nonlinear allocator choice reverses $\Theta(0)$ ’s sign at a reduced path count ($K = 15$); the sign under the correctly specified nonlinear control is not yet established (Section 8.9). No claim in Tables 11–13b should be cited as this paper’s settled conclusion until this is resolved.
- Race-level idiosyncratic uncertainty remains a proxy, not a fitted process (Section 10.1; flagged as a likely-permanent data constraint)
- Opponent reaction is applied to a spending channel (candidate-committee trickle) outside the domain it was estimated on (IE-to-IE reaction) (Section 10.1; flagged as untested)
- The Longstaff–Schwartz solve is not yet folded back into Paper II’s rollout loop as an automatically-applied, continuously-updated reserve policy (Section 10.4; flagged as future work)